

Análisis de Datos y Aprendizaje Máquina con Tensorflow 2.0: Clasificación

2019/09/30

Clasificación K-nearest neighbors y GridSearch (parameter tuning)

Objetivo: Comprender otro tipo de clasificación básica y el ajuste de hiperparámetros.

En la clasificación k-NN, se asigna el elemento a clasificar a la clase donde la mayoría de sus elementos de entrenamiento (vecinos) pertenecen. Si $k = 1$, entonces el objeto simplemente se asigna a la clase de ese vecino más cercano, si $k = 3$, se asigna a la clase donde la mayoría de los tres elementos pertenecen. Este tipo de clasificación es muy común y cuenta con muchas variantes.

<https://scikit-learn.org/stable/modules/neighbors.html#unsupervised-nearest-neighbors>

https://scikit-learn.org/stable/modules/grid_search.html

- Para ajustar el parámetro k con el menor error, así como otros parámetros (parameter tuning), se utiliza GridSearch de sklearn

```
In [1]: import sklearn
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Análisis exploratorio

Etiquetas de clase a valor numérico

- Diagnosis (M = malignant, B = benign)

```
In [2]: %matplotlib inline
df = pd.read_csv("data-breast.csv", index_col=0)
df.head(10)
```

```
Out[2]:
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
id						
842302	M	17.99	10.38	122.80	1001.0	
842517	M	20.57	17.77	132.90	1326.0	
84300903	M	19.69	21.25	130.00	1203.0	
84348301	M	11.42	20.38	77.58	386.1	
84358402	M	20.29	14.34	135.10	1297.0	

843786	M	12.45	15.70	82.57	477.1
844359	M	18.25	19.98	119.60	1040.0
84458202	M	13.71	20.83	90.20	577.9
844981	M	13.00	21.82	87.50	519.8
84501001	M	12.46	24.04	83.97	475.9

	smoothness_mean	compactness_mean	concavity_mean	\
id				
842302	0.11840	0.27760	0.30010	
842517	0.08474	0.07864	0.08690	
84300903	0.10960	0.15990	0.19740	
84348301	0.14250	0.28390	0.24140	
84358402	0.10030	0.13280	0.19800	
843786	0.12780	0.17000	0.15780	
844359	0.09463	0.10900	0.11270	
84458202	0.11890	0.16450	0.09366	
844981	0.12730	0.19320	0.18590	
84501001	0.11860	0.23960	0.22730	

	concave points_mean	symmetry_mean	...	texture_worst	\
id			...		
842302	0.14710	0.2419	...	17.33	
842517	0.07017	0.1812	...	23.41	
84300903	0.12790	0.2069	...	25.53	
84348301	0.10520	0.2597	...	26.50	
84358402	0.10430	0.1809	...	16.67	
843786	0.08089	0.2087	...	23.75	
844359	0.07400	0.1794	...	27.66	
84458202	0.05985	0.2196	...	28.14	
844981	0.09353	0.2350	...	30.73	
84501001	0.08543	0.2030	...	40.68	

	perimeter_worst	area_worst	smoothness_worst	compactness_worst	\
id					
842302	184.60	2019.0	0.1622	0.6656	
842517	158.80	1956.0	0.1238	0.1866	
84300903	152.50	1709.0	0.1444	0.4245	
84348301	98.87	567.7	0.2098	0.8663	
84358402	152.20	1575.0	0.1374	0.2050	
843786	103.40	741.6	0.1791	0.5249	
844359	153.20	1606.0	0.1442	0.2576	
84458202	110.60	897.0	0.1654	0.3682	
844981	106.20	739.3	0.1703	0.5401	
84501001	97.65	711.4	0.1853	1.0580	

	concavity_worst	concave points_worst	symmetry_worst	\
id				
842302	0.7119	0.2654	0.4601	
842517	0.2416	0.1860	0.2750	

84300903	0.4504	0.2430	0.3613
84348301	0.6869	0.2575	0.6638
84358402	0.4000	0.1625	0.2364
843786	0.5355	0.1741	0.3985
844359	0.3784	0.1932	0.3063
84458202	0.2678	0.1556	0.3196
844981	0.5390	0.2060	0.4378
84501001	1.1050	0.2210	0.4366

```

fractal_dimension_worst  Unnamed: 32
id
842302                    0.11890      NaN
842517                    0.08902      NaN
84300903                  0.08758      NaN
84348301                  0.17300      NaN
84358402                  0.07678      NaN
843786                   0.12440      NaN
844359                   0.08368      NaN
84458202                  0.11510      NaN
844981                   0.10720      NaN
84501001                  0.20750      NaN

```

[10 rows x 32 columns]

Análisis Exploratorio

- Verificar valores nulos

In [3]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 569 entries, 842302 to 92751
Data columns (total 32 columns):
diagnosis          569 non-null object
radius_mean        569 non-null float64
texture_mean       569 non-null float64
perimeter_mean     569 non-null float64
area_mean          569 non-null float64
smoothness_mean    569 non-null float64
compactness_mean   569 non-null float64
concavity_mean     569 non-null float64
concave points_mean 569 non-null float64
symmetry_mean      569 non-null float64
fractal_dimension_mean 569 non-null float64
radius_se          569 non-null float64
texture_se         569 non-null float64
perimeter_se       569 non-null float64
area_se            569 non-null float64
smoothness_se      569 non-null float64
compactness_se     569 non-null float64

```

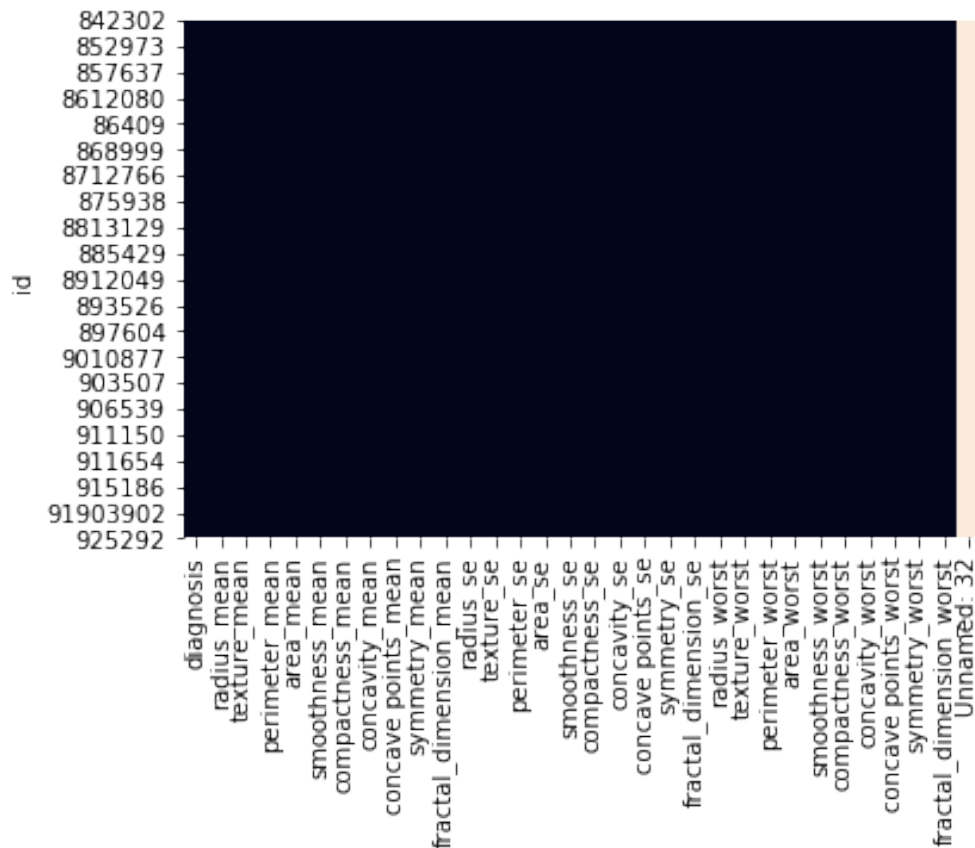
```

concavity_se          569 non-null float64
concave points_se     569 non-null float64
symmetry_se           569 non-null float64
fractal_dimension_se  569 non-null float64
radius_worst          569 non-null float64
texture_worst         569 non-null float64
perimeter_worst       569 non-null float64
area_worst            569 non-null float64
smoothness_worst      569 non-null float64
compactness_worst     569 non-null float64
concavity_worst       569 non-null float64
concave points_worst  569 non-null float64
symmetry_worst        569 non-null float64
fractal_dimension_worst 569 non-null float64
Unnamed: 32           0 non-null float64
dtypes: float64(31), object(1)
memory usage: 146.7+ KB

```

```
In [4]: sns.heatmap(df.isnull(), cbar=False)
```

```
Out[4]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb9e95cc1d0>
```



- Etiquetas a valor numérico

```
In [5]: df = df.replace({'B':0, 'M':1})
df.head()
```

```
Out[5]:
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
id						
842302	1	17.99	10.38	122.80	1001.0	
842517	1	20.57	17.77	132.90	1326.0	
84300903	1	19.69	21.25	130.00	1203.0	
84348301	1	11.42	20.38	77.58	386.1	
84358402	1	20.29	14.34	135.10	1297.0	

	smoothness_mean	compactness_mean	concavity_mean	\
id				
842302	0.11840	0.27760	0.3001	
842517	0.08474	0.07864	0.0869	
84300903	0.10960	0.15990	0.1974	
84348301	0.14250	0.28390	0.2414	
84358402	0.10030	0.13280	0.1980	

	concave points_mean	symmetry_mean	...	texture_worst	\
id			...		
842302	0.14710	0.2419	...	17.33	
842517	0.07017	0.1812	...	23.41	
84300903	0.12790	0.2069	...	25.53	
84348301	0.10520	0.2597	...	26.50	
84358402	0.10430	0.1809	...	16.67	

	perimeter_worst	area_worst	smoothness_worst	compactness_worst	\
id					
842302	184.60	2019.0	0.1622	0.6656	
842517	158.80	1956.0	0.1238	0.1866	
84300903	152.50	1709.0	0.1444	0.4245	
84348301	98.87	567.7	0.2098	0.8663	
84358402	152.20	1575.0	0.1374	0.2050	

	concavity_worst	concave points_worst	symmetry_worst	\
id				
842302	0.7119	0.2654	0.4601	
842517	0.2416	0.1860	0.2750	
84300903	0.4504	0.2430	0.3613	
84348301	0.6869	0.2575	0.6638	
84358402	0.4000	0.1625	0.2364	

	fractal_dimension_worst	Unnamed: 32
id		

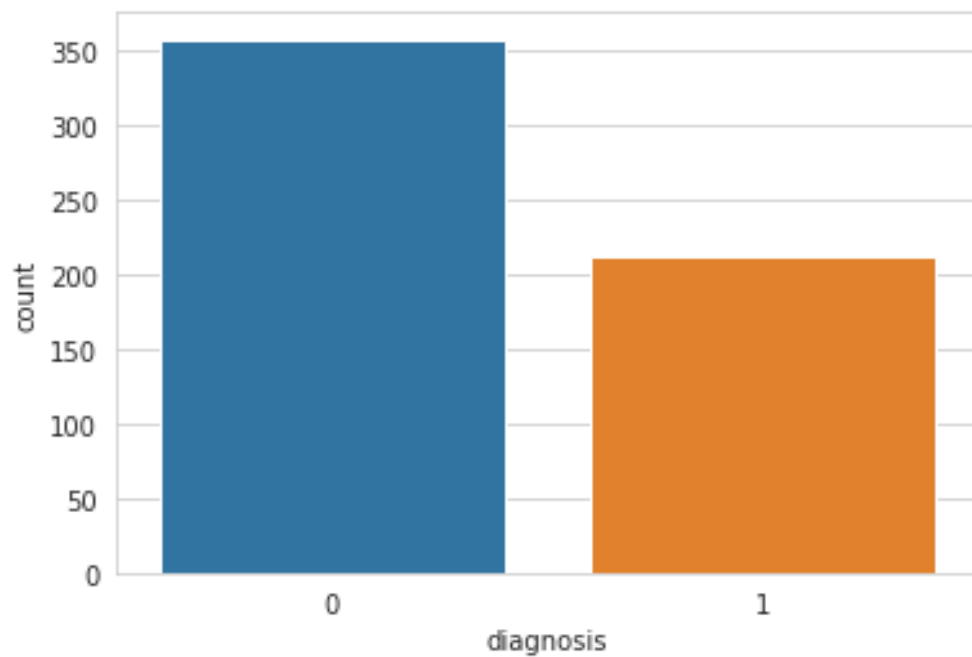
842302	0.11890	NaN
842517	0.08902	NaN
84300903	0.08758	NaN
84348301	0.17300	NaN
84358402	0.07678	NaN

[5 rows x 32 columns]

- Equilibrio de etiquetas

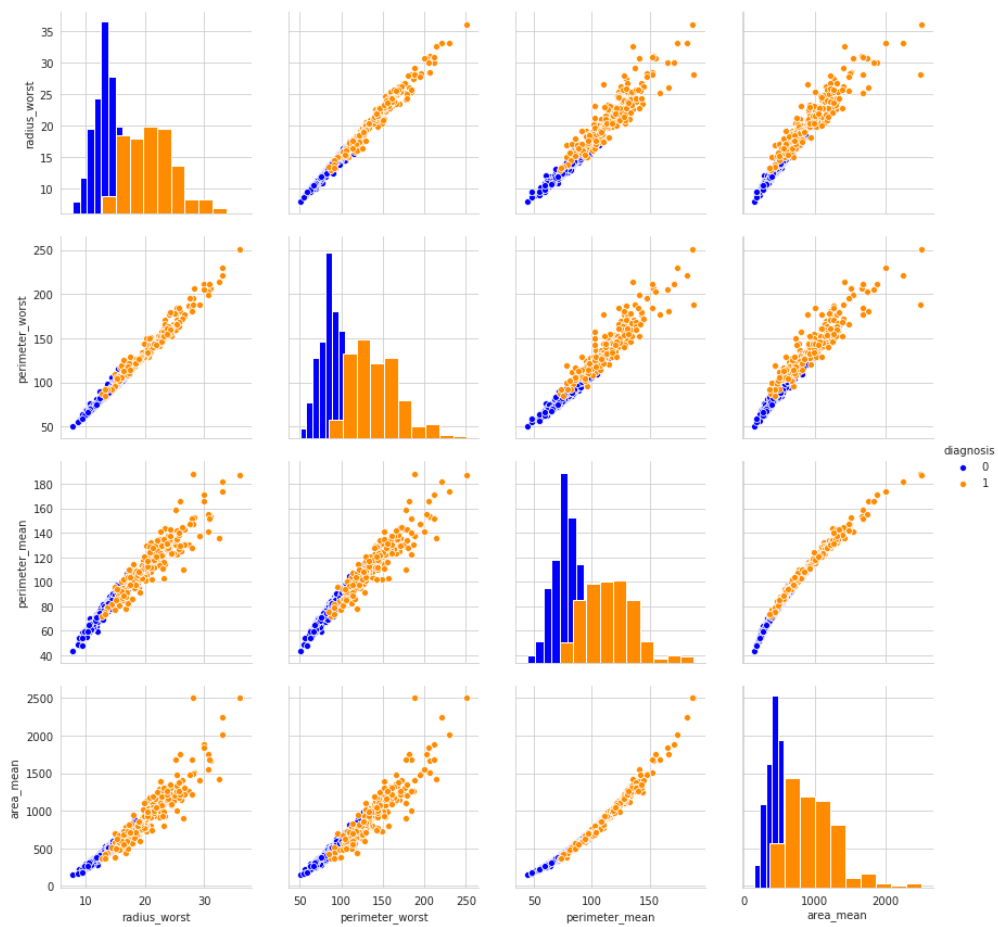
```
In [6]: sns.set_style('whitegrid')
sns.countplot(x='diagnosis',data=df)
```

```
Out[6]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb9e8dee090>
```



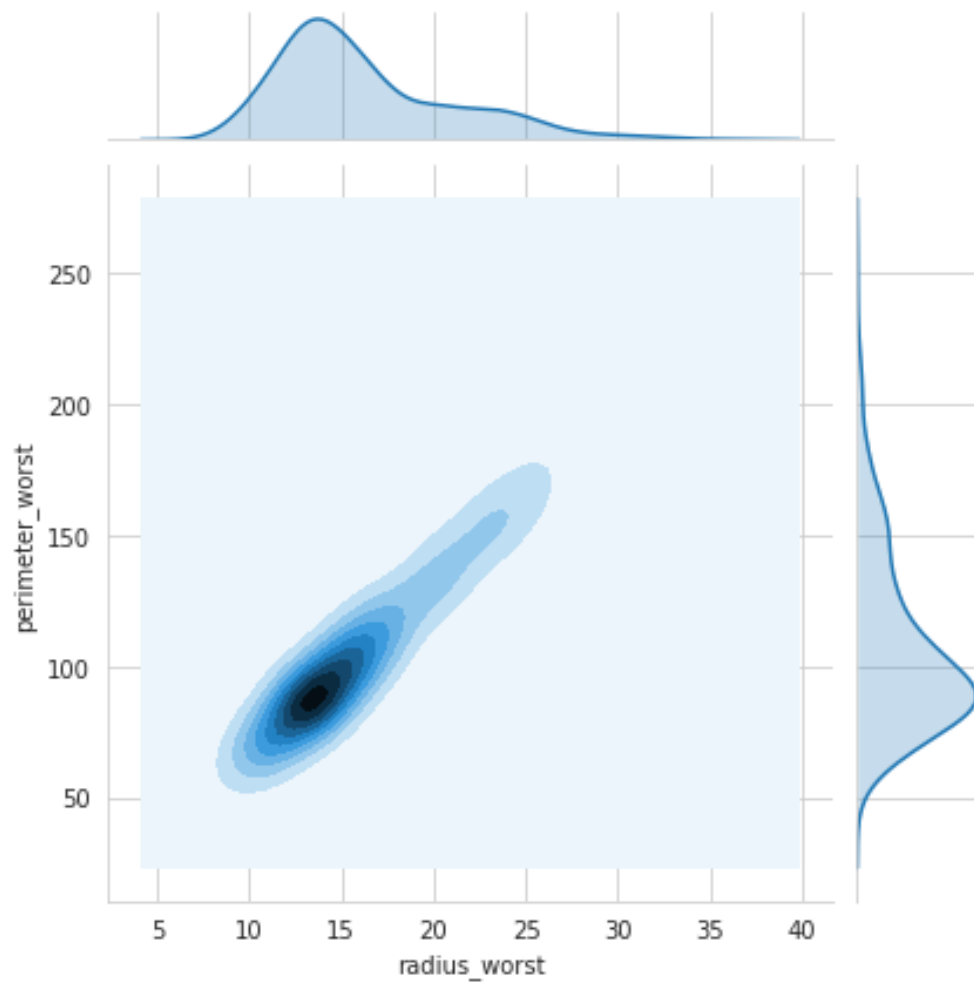
- Distribución de variable

```
In [7]: sns.pairplot(df, vars=["radius_worst", "perimeter_worst", "perimeter_mean", "area_mean"],
palette=sns.color_palette(['blue', 'darkorange']), hue='diagnosis', diag_kind='hist',
plt.show())
```




```
In [8]: sns.jointplot("radius_worst", "perimeter_worst", data=df, kind='kde')
```

```
Out[8]: <seaborn.axisgrid.JointGrid at 0x7fb9e7fa5a50>
```



```
In [9]: sns.jointplot("radius_worst", "perimeter_worst", data=df, color="b").plot_joint(sns.kdeplot)
```

```
Out[9]: <seaborn.axisgrid.JointGrid at 0x7fb9e5ff8110>
```

```
In [10]: from sklearn.neighbors import KNeighborsClassifier
         from sklearn.metrics import accuracy_score
         knn=KNeighborsClassifier()

In [11]: from sklearn.model_selection import train_test_split

         X = df.drop('diagnosis',axis=1)
         X = X.drop('Unnamed: 32',axis=1)
         y = df['diagnosis']
         # dividir datos
         train, test, train_labels, test_labels = train_test_split(X, y,
                                                                    test_size = 0.3, random_state = 22)
```

Ajuste de parámetros de KNN con GridSearchCV

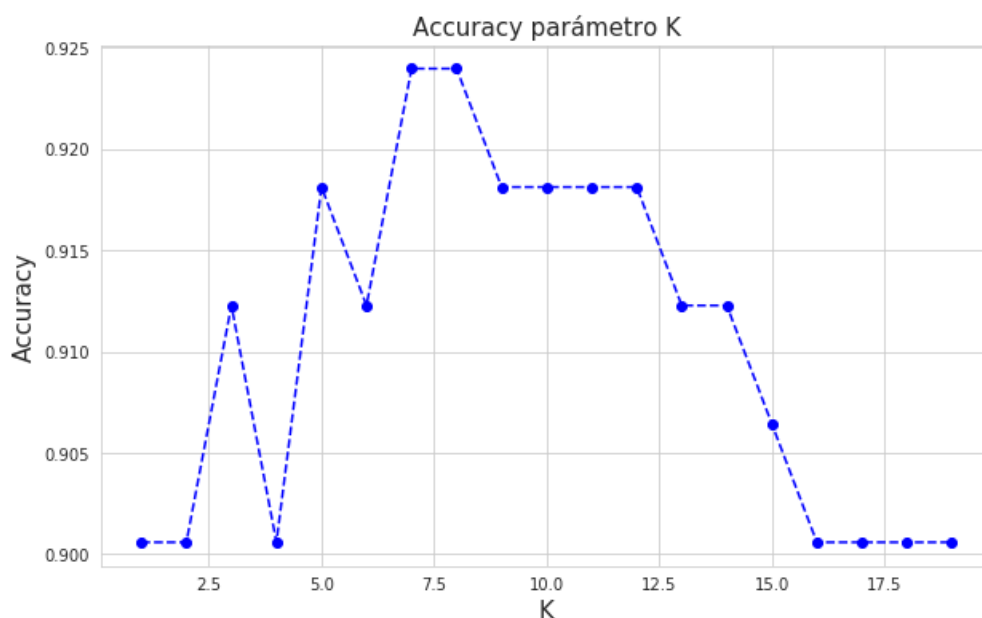
- GridSearchCV realiza una búsqueda por fuerza bruta sobre todos los posibles parámetros

- A diferencia del ejemplo de NaiveBayes, con grid search no se requiere 'for loop' para cada parámetro

```
In [12]: acc = [] #guarda accuracy
         for k in range(1,20):
             knn = KNeighborsClassifier(n_neighbors=k, metric='manhattan')
             knn.fit(train, train_labels)
             pred = knn.predict(test)
             acc.append(accuracy_score(test_labels, pred))

In [13]: plt.figure(figsize=(10,6))
         plt.plot(range(1,20),acc,color='blue', linestyle='--', marker='o')
         plt.title('Accuracy parámetro K', fontsize=15)
         plt.xlabel('K',fontsize=15)
         plt.ylabel('Accuracy',fontsize=15)
```

Out[13]: Text(0, 0.5, 'Accuracy')



- El ciclo indica que k=8 es el mejor parámetro con 'manhattan'

```
In [14]: acc[7]
```

Out[14]: 0.9239766081871345

- Para evaluar los demás parámetros se importa GridSearchCV

```
In [15]: from sklearn.model_selection import GridSearchCV
```

Se evalúan tres parámetros con GridSearchCV

- Además del parámetro 'k' se toma en cuenta 'weights' y 'metric'

```
In [16]: train, test, train_labels, test_labels = train_test_split(X, y,
                                                                    test_size = 0.15, random_state = 22)
```

```
In [17]: import warnings
         warnings.filterwarnings("ignore")
```

```
In [18]: k_vecinos = list(range(1,20))
         weight_options = ["uniform", "distance"]
         metric_options = ["manhattan", 'minkowski']

         param_grid = dict(n_neighbors = k_vecinos, weights = weight_options,
                           metric = metric_options)

         knn = KNeighborsClassifier()

         grid = GridSearchCV(knn, param_grid, scoring = 'accuracy', cv = 10)
         grid.fit(train, train_labels)
```

```
Out[18]: GridSearchCV(cv=10, error_score='raise-deprecating',
                      estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30,
                                                    metric='minkowski',
                                                    metric_params=None, n_jobs=None,
                                                    n_neighbors=5, p=2,
                                                    weights='uniform'),
                      iid='warn', n_jobs=None,
                      param_grid={'metric': ['manhattan', 'minkowski'],
                                   'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12,
                                                  13, 14, 15, 16, 17, 18, 19],
                                   'weights': ['uniform', 'distance']}},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                      scoring='accuracy', verbose=0)
```

Los mejores parámetros

```
In [19]: grid.best_score_
```

```
Out[19]: 0.94824016563147
```

```
In [20]: grid.best_estimator_
```

```
Out[20]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='manhattan',
                              metric_params=None, n_jobs=None, n_neighbors=7, p=2,
                              weights='uniform')
```

```
In [21]: grid.best_params_
```

```
Out[21]: {'metric': 'manhattan', 'n_neighbors': 7, 'weights': 'uniform'}
```

Evaluación de modelo

- Se obtienen las predicciones, informe de clasificación y matriz de confusión.
- Se crea modelo con los mejores parámetros
- Se obtienen las predicciones, informe de clasificación y matriz de confusión.

```
In [22]: predictions = grid.predict(test)
```

```
In [23]: grid.score(test, test_labels)
```

```
Out[23]: 0.9418604651162791
```

```
In [24]: from sklearn.metrics import classification_report, confusion_matrix
print("Predicciones:\n")
print(predictions)
print("\nReporte de clasificación:\n")
print(classification_report(predictions, test_labels))
```

Predicciones:

```
[0 1 1 1 0 0 0 0 0 0 0 0 1 0 0 0 1 1 0 0 1 1 1 0 1 0 1 0 0 0 0 1 0 1 0 1 1
 1 0 1 0 0 0 0 1 0 0 1 0 0 0 1 0 0 1 0 1 0 0 1 0 0 0 0 1 0 0 1 1 0 1 0 1 0
 0 1 0 0 1 0 1 0 1 0 0 1]
```

Reporte de clasificación:

	precision	recall	f1-score	support
0	0.94	0.96	0.95	53
1	0.94	0.91	0.92	33
accuracy			0.94	86
macro avg	0.94	0.94	0.94	86
weighted avg	0.94	0.94	0.94	86

```
In [25]: print("Confusion matrix")
conf_mat=confusion_matrix(predictions, test_labels)
print(conf_mat)
```

Confusion matrix

```
[[51  2]
 [ 3 30]]
```

- Ajustando los parámetros se obtiene un modelo mas potente
- Probar los clasificadores aprendidos con un diferente dataset
- Experimentar con parámetros y GridSearch