

Análisis de Datos y Aprendizaje Máquina con Tensorflow 2.0: Perceptrón Multicapa

2019/09/30

1 Optimizadores

Objetivo: Conocer los diferentes optimizadores para entrenar redes neuronales

- Referencia Adam: <https://arxiv.org/abs/1412.6980>
- Referencia Adadelta: <https://arxiv.org/abs/1212.5701>

SGD es un algoritmo que se emplea para minimizar una función objetivo, respecto a algunos parámetros. Este proceso es iterativo, tomando en cuenta un parámetro de learning rate, que establece que tanto se actualizan los parámetros que consiguen los resultados deseados.

- Una variante de SGD efectiva es Adam (Estimación de Momentum Adaptativo), también se emplean RMSProp, Adadelta y Adagrad.

1.0.1 Se itera para encontrar el valor mínimo de una función utilizando la dirección del gradiente

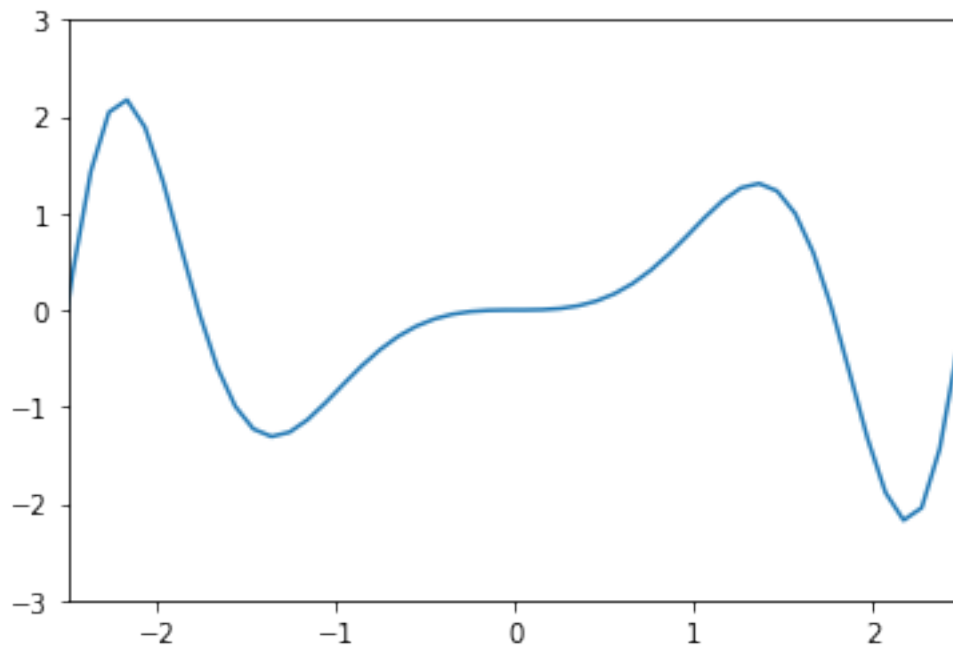
- Se busca el valor mínimo de $J(x) = x \sin(x^2)$

```
In [28]: import numpy as np
import matplotlib.pyplot as plt
x = np.linspace(-5, 5, 100)

def fn(x):# funcion
    return x*np.sin(x*x)

y = list(map(fn, x))

plt.plot(x, y)
plt.ylim(-3, 3)
plt.xlim(-2.5, 2.5)
plt.show()
```



- Se selecciona el número de iteraciones = 10 y el learning rate = 0.01
- La posición inicial se denota con 'ini'

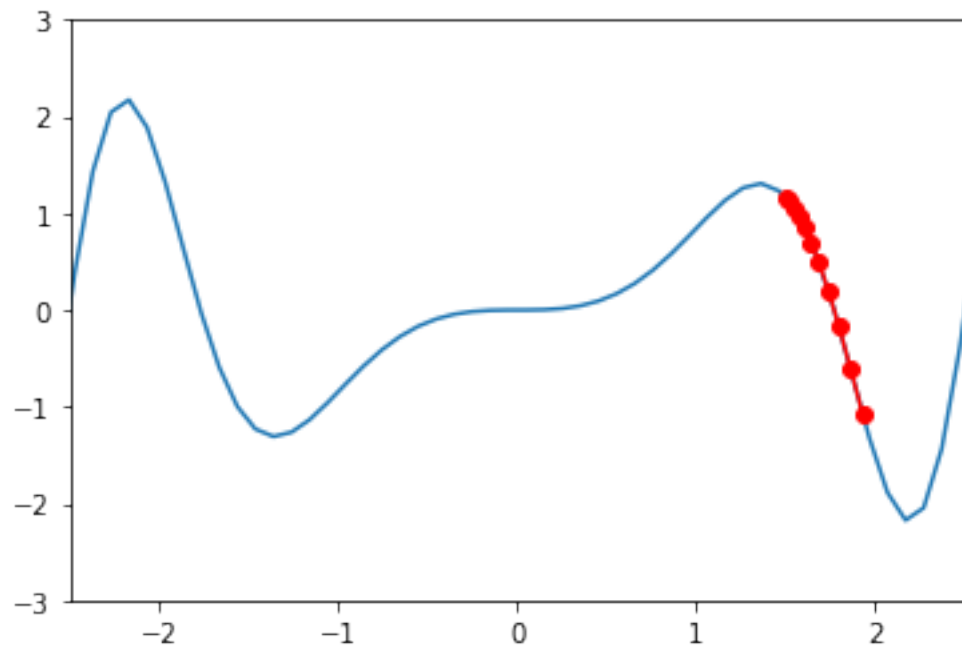
```
In [2]: def dif(x): ##derivada de función
        return (np.sin(x*x) + 2*x*x*np.cos(x*x))
```

```
In [3]: def history(ini, lr, it = 10):
        history = []
        history.append(ini)
        x = ini
        for _ in range(0, it):# n iteraciones
            x = x - lr*dif(x)
            history.append(x)
        return history
```

```
In [4]: h=history(1.5,0.01,10)
```

```
In [5]: x = np.linspace(-5, 5, 100)
        yh = list(map(fn, h))

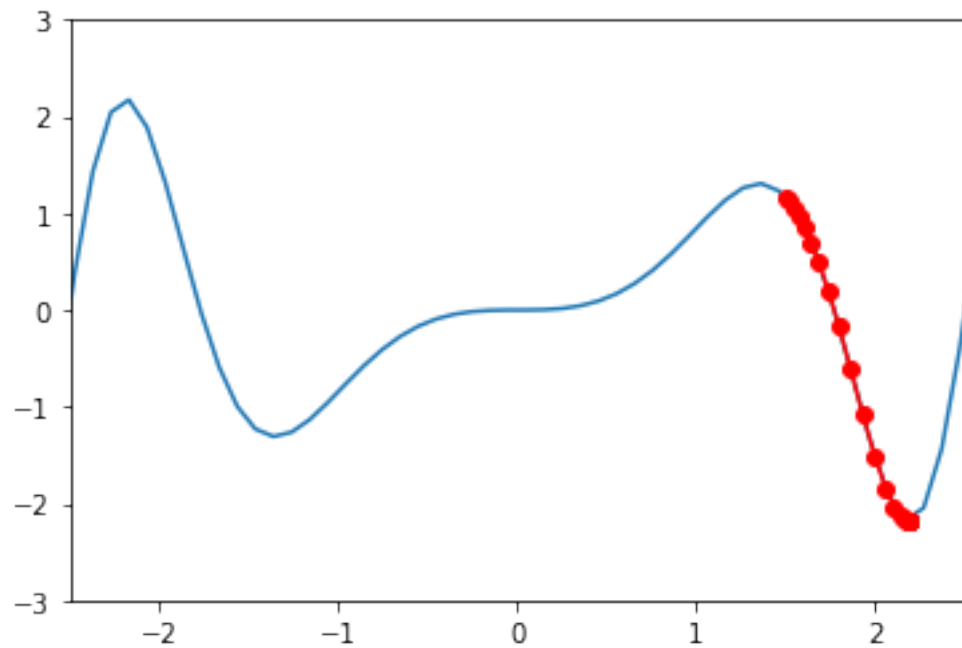
        plt.plot(x, y, '-')
        plt.plot(h, yh, 'o-', c='red')
        plt.ylim(-3, 3)
        plt.xlim(-2.5, 2.5)
        plt.show()
```



- Aumentando iteraciones a 100

```
In [6]: h=history(1.5,0.01,100)
        x = np.linspace(-5, 5, 100)
        yh = list(map(fn, h))

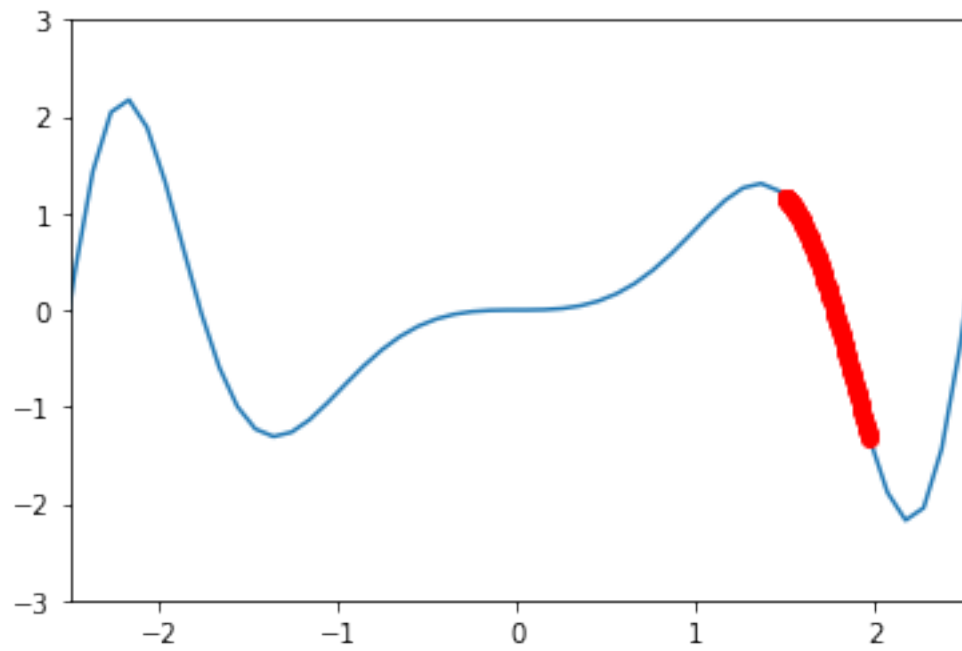
        plt.plot(x, y, '-')
        plt.plot(h, yh, 'o-', c='red')
        plt.ylim(-3, 3)
        plt.xlim(-2.5, 2.5)
        plt.show()
```



- Un learning rate bajo puede requerir mayor número de iteraciones

```
In [7]: h=history(1.5,0.001,100)# 100 iteraciones
        x = np.linspace(-5, 5, 100)
        yh = list(map(fn, h))

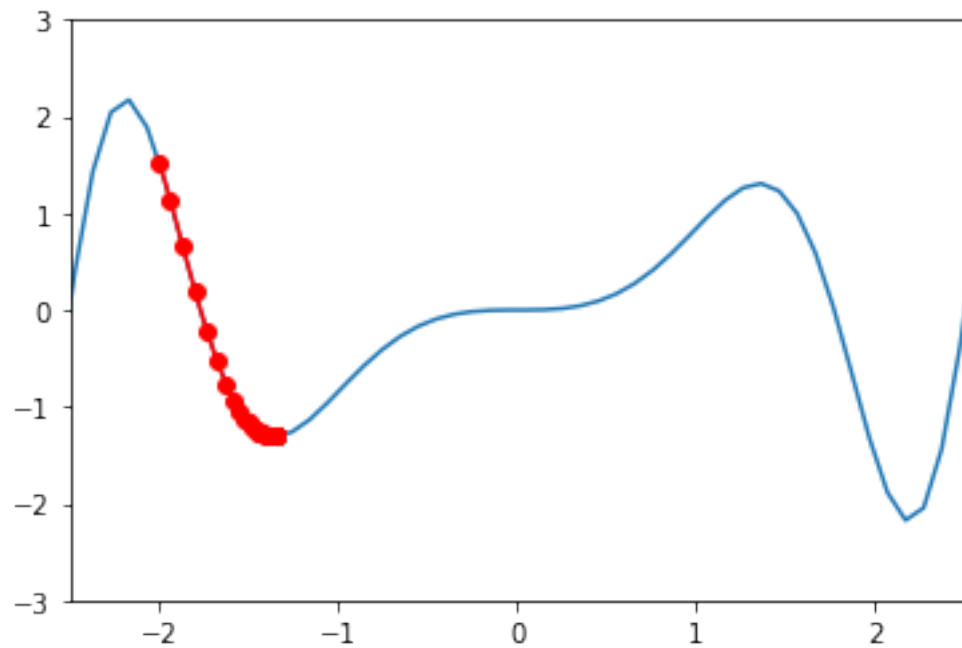
        plt.plot(x, y, '-')
        plt.plot(h, yh,'o-', c='red')
        plt.ylim(-3, 3)
        plt.xlim(-2.5, 2.5)
        plt.show()
```



- Diferente punto de inicio

```
In [8]: h=history(-2,0.01,100)# 100 iteraciones
        x = np.linspace(-5, 5, 100)
        yh = list(map(fn, h))

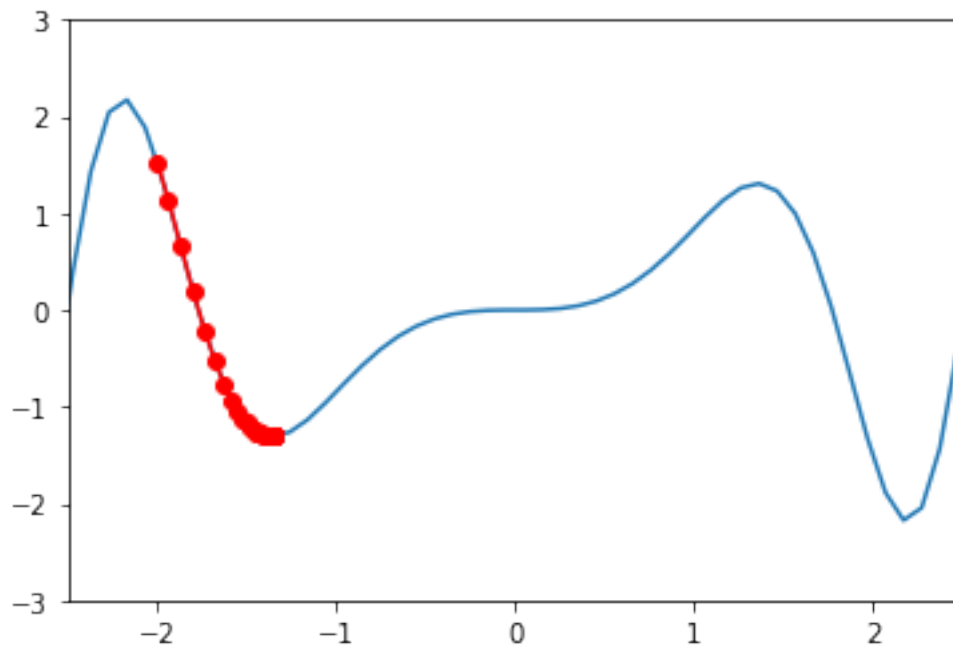
        plt.plot(x, y, '-')
        plt.plot(h, yh,'o-', c='red')
        plt.ylim(-3, 3)
        plt.xlim(-2.5, 2.5)
        plt.show()
```



- Más iteraciones

```
In [9]: h=history(-2,0.01,1000)# 1000 iteraciones
        x = np.linspace(-5, 5, 100)
        yh = list(map(fn, h))

        plt.plot(x, y, '-')
        plt.plot(h, yh, 'o-', c='red')
        plt.ylim(-3, 3)
        plt.xlim(-2.5, 2.5)
        plt.show()
```



- La función no puede encontrar el mínimo, es por eso que existen variantes del algoritmo

2 Tensorflow optimizadores

- Tensorflow implementa varios optimizadores
- Se clasificarán imágenes. Los pixeles son la entrada de la red neuronal. Al finalizar el entrenamiento, la red neuronal habrá aprendido a reconocer dígitos.

```
In [10]: import tensorflow as tf
         from tensorflow import keras
         from tensorflow.keras.models import Sequential
         from tensorflow.keras.layers import Dense
         from tensorflow.keras import backend as K
         K.clear_session()

         mnist = keras.datasets.mnist

         (x_train, y_train), (x_test, y_test) = mnist.load_data()


In [11]: print(x_train.shape)
         print(y_train.shape)
         print(x_test.shape)
         print(y_test.shape)
```

```
(60000, 28, 28)
(60000,)
(10000, 28, 28)
(10000,)
```

2.1 Leer Dataset

```
In [12]: fig, ax = plt.subplots(nrows=2, ncols=5, sharex=True, sharey=True,)
        ax = ax.flatten()
        for i in range(10):
            img = x_train[y_train == i][0].reshape(28, 28)
            ax[i].imshow(img, cmap='Greys', interpolation='nearest')

        ax[0].set_xticks([])
        ax[0].set_yticks([])
        plt.tight_layout()
        plt.show()
```



- Se modifica la forma de los datos de 2-d (n, 28, 28) a 1-d (n, 784)

```
In [13]: x_train = x_train.reshape(x_train.shape[0], -1)
        x_test = x_test.reshape(x_test.shape[0], -1)
```

```
print(x_train.shape) # (60000, 784)
print(y_train.shape) # (60000,)
print(x_test.shape)  # (10000, 784)
print(y_test.shape)  # (10000,)
```

```
(60000, 784)
(60000,)
```



```
(10000, 784)
(10000,)
```

```
In [14]: epoch = 70
         verbose = 0
         batch = 50
```

2.2 Adam

```
In [15]: def make_model():
         model = Sequential()

         model.add(Dense(40, input_shape = (784, ), activation = 'relu'))
         model.add(Dense(40, activation = 'relu'))
         model.add(Dense(40, activation = 'relu'))
         model.add(Dense(10, activation = 'softmax'))

         model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
                       metrics=['accuracy'])
         return model
```

```
In [16]: model = make_model()
```

```
         model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 40)	31400
dense_1 (Dense)	(None, 40)	1640
dense_2 (Dense)	(None, 40)	1640
dense_3 (Dense)	(None, 10)	410

Total params: 35,090

Trainable params: 35,090

Non-trainable params: 0

```
In [17]: history1 = model.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = epoch, verbose = verbose)
```

```
In [18]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
```

```
         print('\nTest accuracy:', test_acc)
```

10000/1 - 0s - loss: 0.1586 - accuracy: 0.9633

Test accuracy: 0.9633

2.3 SGD

```
In [19]: def make_model():
         model = Sequential()

         model.add(Dense(40, input_shape = (784, ), activation = 'relu'))
         model.add(Dense(40, activation = 'relu'))
         model.add(Dense(40, activation = 'relu'))
         model.add(Dense(10, activation = 'softmax'))

         model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy',
                       metrics=['accuracy'])
         return model
```

```
In [20]: model = make_model()
```

```
         model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 40)	31400
dense_5 (Dense)	(None, 40)	1640
dense_6 (Dense)	(None, 40)	1640
dense_7 (Dense)	(None, 10)	410

Total params: 35,090
Trainable params: 35,090
Non-trainable params: 0

```
In [21]: history2 = model.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = epoch, verbose = verbose)
```

```
In [22]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
```

```
         print('\nTest accuracy:', test_acc)
```

10000/1 - 0s - loss: 2.3024 - accuracy: 0.1135

Test accuracy: 0.1135

2.4 Adagrad

```
In [23]: def make_model():
         model = Sequential()

         model.add(Dense(40, input_shape = (784, ), activation = 'relu'))
         model.add(Dense(40, activation = 'relu'))
         model.add(Dense(40, activation = 'relu'))
         model.add(Dense(10, activation = 'softmax'))

         model.compile(optimizer='adagrad', loss='sparse_categorical_crossentropy',
                       metrics=['accuracy'])
         return model
```

```
In [24]: model = make_model()
```

```
model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 40)	31400
dense_9 (Dense)	(None, 40)	1640
dense_10 (Dense)	(None, 40)	1640
dense_11 (Dense)	(None, 10)	410

=====
Total params: 35,090
Trainable params: 35,090
Non-trainable params: 0
=====

```
In [25]: history3 = model.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = epoch, verbose = verbose)
```

```
In [26]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
```

```
print('\nTest accuracy:', test_acc)
```

10000/1 - 1s - loss: 0.4100 - accuracy: 0.8699

Test accuracy: 0.8699

```

In [27]: #plot
plt.figure(figsize=(10,9))
plt.plot(history1.history['accuracy'])
plt.plot(history1.history['val_accuracy'])

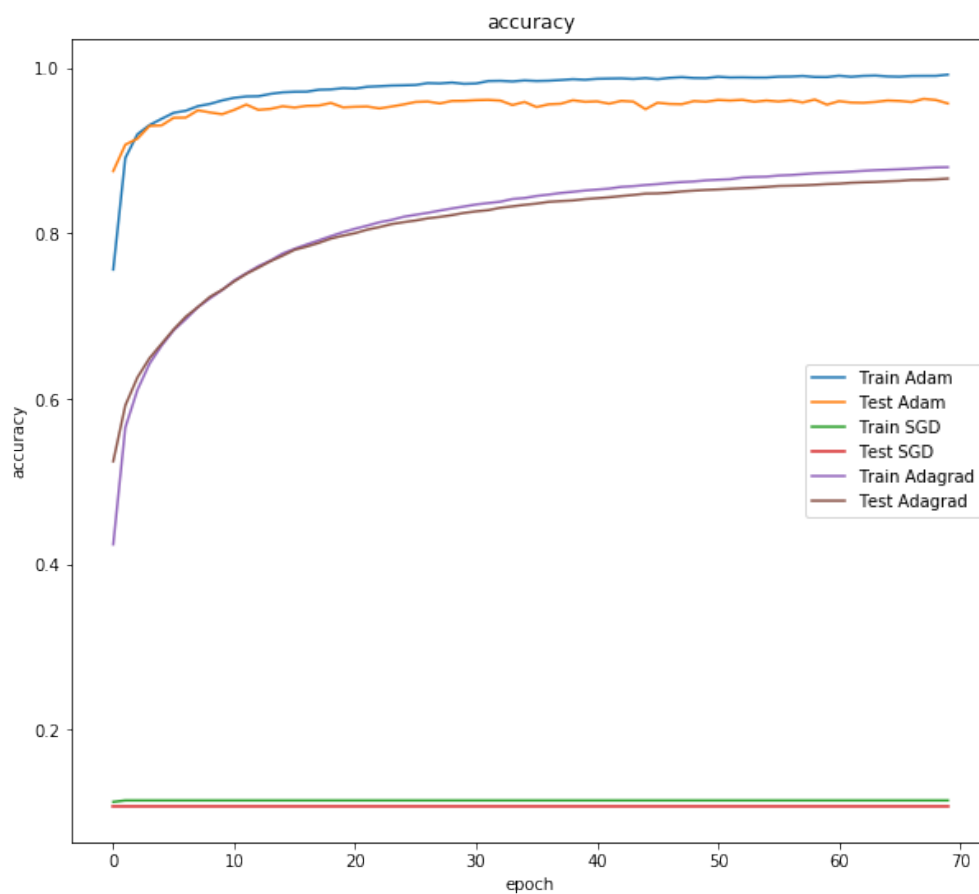
plt.plot(history2.history['accuracy'])
plt.plot(history2.history['val_accuracy'])

plt.plot(history3.history['accuracy'])
plt.plot(history3.history['val_accuracy'])
plt.legend(['Train Adam', 'Test Adam',
            'Train SGD', 'Test SGD',
            'Train Adagrad', 'Test Adagrad'])

plt.title('accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')

plt.show()

```



- Experimentar con diferentes parámetros y optimizadores para obtener mejores resultados y menor tiempo de entrenamiento