

Análisis de Datos y Aprendizaje Máquina con Tensorflow 2.0: Redes neuronales recurrentes

2019/09/30

RNN - Clasificación de Texto


- Objetivo: Implementar RNN y LSTM para la clasificación de secuencias (texto). Aplicar inicialización de pesos y dropout en modelos recurrentes.
- Conocer efectos de inicialización y dropout en RNN y comparar con MLP

Redes Neuronales Recurrentes

- Cada capa en cada iteración comparte parámetros
- 'sparse_categorical_crossentropy' se utiliza para varias clases

Leer Dataset

```
In [1]: from tensorflow.keras.models import Model
        from tensorflow.keras.layers import LSTM, Embedding, Dense, SimpleRNN
        from tensorflow.keras.datasets import reuters
        from tensorflow.keras.models import Sequential
        from tensorflow import keras
        import numpy as np
        from tensorflow.keras.preprocessing.sequence import pad_sequences
        import matplotlib.pyplot as plt
```



RNN

11,228 noticias de Reuters, con más de 46 temas.

```
In [2]: # parámetros
        emb_dim = 64
        num_words = 8000
        max_words = 40
```

```
In [3]: (x_train, y_train), (x_test, y_test) = reuters.load_data(num_words = num_words, maxlen=ma
```

```
In [4]: x_train.shape
```

```
Out[4]: (1084,)
```

```
In [5]: print('Noticia')
        print(x_train[0])
        print('Etiqueta')
        print(y_train[0])
```

Noticia

[1, 245, 273, 207, 156, 53, 74, 160, 26, 14, 46, 296, 26, 39, 74, 2979, 3554, 14, 46, 4689, 4329,

Etiqueta

3

Palabras de noticia

```
In [6]: wordDict = {y:x for x,y in reuters.get_word_index().items()}
        res = []
        for index in x_train[0]:
            res.append(wordDict.get(index - 3))
        print('Noticia: ',res,'Longitud noticia: ', len(res))
```

Noticia: [None, 'period', 'ended', 'december', '31', 'shr', 'profit', '11', 'cts', 'vs', 'loss',

```
In [7]: x_train = pad_sequences(x_train, maxlen=max_words, padding = 'post')
        x_test = pad_sequences(x_test, maxlen=max_words, padding = 'post')
```

```
In [8]: epoch = 40
        verbose = 1
        batch = 50
```

```
In [9]: print(x_train.shape)
        print(x_test.shape)
        print(y_train.shape)
        print(y_test.shape)
```

(1084, 40)

(272, 40)

(1084,)

(272,)

```
In [10]: x_train[0]
```

```
Out[10]: array([[ 1, 245, 273, 207, 156, 53, 74, 160, 26, 14, 46,
                296, 26, 39, 74, 2979, 3554, 14, 46, 4689, 4329, 86,
                61, 3499, 4795, 14, 61, 451, 4329, 17, 12, 0, 0,
                0, 0, 0, 0, 0, 0, 0], dtype=int32)
```

Deep RNN

- Cuando se conectan varias capas de RNNs se modifica el parámetro 'return_sequences'
- Se inicializan los pesos con 'glorot_uniform'

```
In [11]: def deep_rnn():
```

```
    model = Sequential()
    model.add(Embedding(num_words, emb_dim))
    model.add(SimpleRNN(32, return_sequences = True, recurrent_initializer='glorot_uniform'))
    model.add(SimpleRNN(32, return_sequences = False, recurrent_initializer='glorot_uniform'))
    model.add(Dense(46, activation = 'softmax'))
```

```
    model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
    return model
```

```
In [12]: model = deep_rnn()
```

```
    model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, None, 64)	512000
simple_rnn (SimpleRNN)	(None, None, 32)	3104
simple_rnn_1 (SimpleRNN)	(None, 32)	2080
dense (Dense)	(None, 46)	1518

=====
Total params: 518,702
Trainable params: 518,702
Non-trainable params: 0
=====

```
In [13]: history = model.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = epoch, verbose = verbose)
```

```
Train on 758 samples, validate on 326 samples
```

```
Epoch 1/40
```

```
758/758 [=====] - 2s 3ms/sample - loss: 2.7716 - accuracy: 0.6240 - val_loss: 2.7716 - val_accuracy: 0.6240
```

```
Epoch 2/40
```

```

758/758 [=====] - 1s 886us/sample - loss: 1.5105 - accuracy: 0.7718 - va
Epoch 3/40
758/758 [=====] - 1s 998us/sample - loss: 1.1607 - accuracy: 0.7718 - va
Epoch 4/40
758/758 [=====] - 1s 937us/sample - loss: 1.0992 - accuracy: 0.7718 - va
Epoch 5/40
758/758 [=====] - 1s 1ms/sample - loss: 1.0718 - accuracy: 0.7718 - val_
Epoch 6/40
758/758 [=====] - 1s 824us/sample - loss: 1.0595 - accuracy: 0.7718 - va
Epoch 7/40
758/758 [=====] - 1s 922us/sample - loss: 1.0503 - accuracy: 0.7718 - va
Epoch 8/40
758/758 [=====] - 1s 877us/sample - loss: 1.0444 - accuracy: 0.7718 - va
Epoch 9/40
758/758 [=====] - 1s 855us/sample - loss: 1.0403 - accuracy: 0.7718 - va
Epoch 10/40
758/758 [=====] - 1s 870us/sample - loss: 1.0373 - accuracy: 0.7718 - va
Epoch 11/40
758/758 [=====] - 1s 870us/sample - loss: 1.0345 - accuracy: 0.7718 - va
Epoch 12/40
758/758 [=====] - 1s 890us/sample - loss: 1.0320 - accuracy: 0.7718 - va
Epoch 13/40
758/758 [=====] - 1s 1ms/sample - loss: 1.0313 - accuracy: 0.7718 - val_
Epoch 14/40
758/758 [=====] - 1s 846us/sample - loss: 1.0304 - accuracy: 0.7718 - va
Epoch 15/40
758/758 [=====] - 1s 809us/sample - loss: 1.0293 - accuracy: 0.7718 - va
Epoch 16/40
758/758 [=====] - 1s 900us/sample - loss: 1.0280 - accuracy: 0.7718 - va
Epoch 17/40
758/758 [=====] - 1s 812us/sample - loss: 1.0279 - accuracy: 0.7718 - va
Epoch 18/40
758/758 [=====] - 1s 879us/sample - loss: 1.0261 - accuracy: 0.7718 - va
Epoch 19/40
758/758 [=====] - 1s 870us/sample - loss: 1.0267 - accuracy: 0.7718 - va
Epoch 20/40
758/758 [=====] - 1s 817us/sample - loss: 1.0268 - accuracy: 0.7718 - va
Epoch 21/40
758/758 [=====] - 1s 986us/sample - loss: 1.0250 - accuracy: 0.7718 - va
Epoch 22/40
758/758 [=====] - 1s 915us/sample - loss: 1.0246 - accuracy: 0.7718 - va
Epoch 23/40
758/758 [=====] - 1s 830us/sample - loss: 1.0241 - accuracy: 0.7718 - va
Epoch 24/40
758/758 [=====] - 1s 882us/sample - loss: 1.0237 - accuracy: 0.7718 - va
Epoch 25/40
758/758 [=====] - 1s 880us/sample - loss: 1.0230 - accuracy: 0.7718 - va
Epoch 26/40
758/758 [=====] - 1s 866us/sample - loss: 1.0235 - accuracy: 0.7718 - va

```

```

Epoch 27/40
758/758 [=====] - 1s 939us/sample - loss: 1.0239 - accuracy: 0.7718 - va
Epoch 28/40
758/758 [=====] - 1s 958us/sample - loss: 1.0244 - accuracy: 0.7718 - va
Epoch 29/40
758/758 [=====] - 1s 933us/sample - loss: 1.0221 - accuracy: 0.7718 - va
Epoch 30/40
758/758 [=====] - 1s 858us/sample - loss: 1.0229 - accuracy: 0.7718 - va
Epoch 31/40
758/758 [=====] - 1s 838us/sample - loss: 1.0228 - accuracy: 0.7718 - va
Epoch 32/40
758/758 [=====] - 1s 922us/sample - loss: 1.0225 - accuracy: 0.7718 - va
Epoch 33/40
758/758 [=====] - 1s 940us/sample - loss: 1.0227 - accuracy: 0.7718 - va
Epoch 34/40
758/758 [=====] - 1s 923us/sample - loss: 1.0220 - accuracy: 0.7718 - va
Epoch 35/40
758/758 [=====] - 1s 943us/sample - loss: 1.0213 - accuracy: 0.7718 - va
Epoch 36/40
758/758 [=====] - 1s 1ms/sample - loss: 1.0220 - accuracy: 0.7718 - val_
Epoch 37/40
758/758 [=====] - 1s 884us/sample - loss: 1.0227 - accuracy: 0.7718 - va
Epoch 38/40
758/758 [=====] - 1s 975us/sample - loss: 1.0214 - accuracy: 0.7718 - va
Epoch 39/40
758/758 [=====] - 1s 967us/sample - loss: 1.0215 - accuracy: 0.7718 - va
Epoch 40/40
758/758 [=====] - 1s 1ms/sample - loss: 1.0213 - accuracy: 0.7718 - val_

```

```
In [14]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
```

```
print('\nTest accuracy:', test_acc)
```

```
272/1 - 0s - loss: 1.1919 - accuracy: 0.8272
```

```
Test accuracy: 0.8272059
```


```
In [15]: #plot
plt.figure(figsize=(10,9))

plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('accuracy')
plt.legend(['train', 'test'])
plt.grid()
plt.subplot(212)
plt.plot(history.history['loss'])

```

```
plt.plot(history.history['val_loss'])
plt.title('loss')
plt.legend(['train', 'test'])
plt.grid()
```

```
plt.show()
```



LSTM

- Desempeño de LSTM con una capa vs. deep RNN

```
In [16]: def lstm():
          model = Sequential()
          model.add(Embedding(num_words, emb_dim))
          model.add(LSTM(32, return_sequences = False))
```

```

        model.add(Dense(46, activation = 'softmax'))

        model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

        return model

```

```

In [17]: model = lstm()
        model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, None, 64)	512000
lstm (LSTM)	(None, 32)	12416
dense_1 (Dense)	(None, 46)	1518
Total params: 525,934		
Trainable params: 525,934		
Non-trainable params: 0		

```

In [18]: history = model.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = epoch, verbose = verbose)

```

Train on 758 samples, validate on 326 samples

```

Epoch 1/40
758/758 [=====] - 2s 3ms/sample - loss: 3.7126 - accuracy: 0.5858 - val_loss: 3.7126
Epoch 2/40
758/758 [=====] - 0s 257us/sample - loss: 2.7176 - accuracy: 0.7718 - val_loss: 2.7176
Epoch 3/40
758/758 [=====] - 0s 256us/sample - loss: 1.5607 - accuracy: 0.7718 - val_loss: 1.5607
Epoch 4/40
758/758 [=====] - 0s 251us/sample - loss: 1.1844 - accuracy: 0.7718 - val_loss: 1.1844
Epoch 5/40
758/758 [=====] - 0s 255us/sample - loss: 1.0998 - accuracy: 0.7718 - val_loss: 1.0998
Epoch 6/40
758/758 [=====] - 0s 245us/sample - loss: 1.0740 - accuracy: 0.7718 - val_loss: 1.0740
Epoch 7/40
758/758 [=====] - 0s 231us/sample - loss: 1.0598 - accuracy: 0.7718 - val_loss: 1.0598
Epoch 8/40
758/758 [=====] - 0s 243us/sample - loss: 1.0517 - accuracy: 0.7718 - val_loss: 1.0517
Epoch 9/40
758/758 [=====] - 0s 252us/sample - loss: 1.0458 - accuracy: 0.7718 - val_loss: 1.0458
Epoch 10/40
758/758 [=====] - 0s 259us/sample - loss: 1.0410 - accuracy: 0.7718 - val_loss: 1.0410
Epoch 11/40

```

758/758 [=====] - 0s 262us/sample - loss: 1.0380 - accuracy: 0.7718 - va
 Epoch 12/40
 758/758 [=====] - 0s 244us/sample - loss: 1.0353 - accuracy: 0.7718 - va
 Epoch 13/40
 758/758 [=====] - 0s 352us/sample - loss: 1.0332 - accuracy: 0.7718 - va
 Epoch 14/40
 758/758 [=====] - 0s 307us/sample - loss: 1.0320 - accuracy: 0.7718 - va
 Epoch 15/40
 758/758 [=====] - 0s 299us/sample - loss: 1.0312 - accuracy: 0.7718 - va
 Epoch 16/40
 758/758 [=====] - 0s 266us/sample - loss: 1.0291 - accuracy: 0.7718 - va
 Epoch 17/40
 758/758 [=====] - 0s 275us/sample - loss: 1.0282 - accuracy: 0.7718 - va
 Epoch 18/40
 758/758 [=====] - 0s 352us/sample - loss: 1.0284 - accuracy: 0.7718 - va
 Epoch 19/40
 758/758 [=====] - 0s 301us/sample - loss: 1.0262 - accuracy: 0.7718 - va
 Epoch 20/40
 758/758 [=====] - 0s 253us/sample - loss: 1.0268 - accuracy: 0.7718 - va
 Epoch 21/40
 758/758 [=====] - 0s 312us/sample - loss: 1.0253 - accuracy: 0.7718 - va
 Epoch 22/40
 758/758 [=====] - 0s 328us/sample - loss: 1.0250 - accuracy: 0.7718 - va
 Epoch 23/40
 758/758 [=====] - 0s 264us/sample - loss: 1.0248 - accuracy: 0.7718 - va
 Epoch 24/40
 758/758 [=====] - 0s 274us/sample - loss: 1.0243 - accuracy: 0.7718 - va
 Epoch 25/40
 758/758 [=====] - 0s 285us/sample - loss: 1.0233 - accuracy: 0.7718 - va
 Epoch 26/40
 758/758 [=====] - 0s 293us/sample - loss: 1.0251 - accuracy: 0.7718 - va
 Epoch 27/40
 758/758 [=====] - 0s 272us/sample - loss: 1.0225 - accuracy: 0.7718 - va
 Epoch 28/40
 758/758 [=====] - 0s 277us/sample - loss: 1.0226 - accuracy: 0.7718 - va
 Epoch 29/40
 758/758 [=====] - 0s 263us/sample - loss: 1.0229 - accuracy: 0.7718 - va
 Epoch 30/40
 758/758 [=====] - 0s 268us/sample - loss: 1.0225 - accuracy: 0.7718 - va
 Epoch 31/40
 758/758 [=====] - 0s 301us/sample - loss: 1.0224 - accuracy: 0.7718 - va
 Epoch 32/40
 758/758 [=====] - 0s 270us/sample - loss: 1.0218 - accuracy: 0.7718 - va
 Epoch 33/40
 758/758 [=====] - 0s 307us/sample - loss: 1.0229 - accuracy: 0.7718 - va
 Epoch 34/40
 758/758 [=====] - 0s 286us/sample - loss: 1.0221 - accuracy: 0.7718 - va
 Epoch 35/40
 758/758 [=====] - 0s 296us/sample - loss: 1.0223 - accuracy: 0.7718 - va


```

Epoch 36/40
758/758 [=====] - 0s 258us/sample - loss: 1.0226 - accuracy: 0.7718 - va
Epoch 37/40
758/758 [=====] - 0s 292us/sample - loss: 1.0210 - accuracy: 0.7718 - va
Epoch 38/40
758/758 [=====] - 0s 244us/sample - loss: 1.0211 - accuracy: 0.7718 - va
Epoch 39/40
758/758 [=====] - 0s 265us/sample - loss: 1.0199 - accuracy: 0.7718 - va
Epoch 40/40
758/758 [=====] - 0s 273us/sample - loss: 1.0198 - accuracy: 0.7718 - va

```

```
In [19]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
```

```
print('\nTest accuracy:', test_acc)
```


```
272/1 - 0s - loss: 1.1993 - accuracy: 0.8272
```

```
Test accuracy: 0.8272059
```

```
In [20]: #plot
plt.figure(figsize=(10,9))

plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('accuracy')
plt.legend(['train', 'test'])
plt.grid()
plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('loss')
plt.legend(['train', 'test'])
plt.grid()

plt.show()
```



Deep LSTM

- LSTM cuentan inicializador 'orthogonal' por defecto

```
In [21]: def deep_lstm():
        model = Sequential()
        model.add(Embedding(num_words, emb_dim))
        model.add(LSTM(32, return_sequences = True, recurrent_initializer='orthogonal'))
        model.add(LSTM(32, return_sequences = False, recurrent_initializer='orthogonal'))
        model.add(Dense(46, activation = 'softmax'))
        model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

        return model
```

```
In [22]: model = deep_lstm()
        model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, None, 64)	512000
lstm_1 (LSTM)	(None, None, 32)	12416
lstm_2 (LSTM)	(None, 32)	8320
dense_2 (Dense)	(None, 46)	1518

Total params: 534,254
 Trainable params: 534,254
 Non-trainable params: 0

```
In [23]: history1 = model.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = epoch, verbose = verbose)
```

Train on 758 samples, validate on 326 samples

```
Epoch 1/40
758/758 [=====] - 3s 4ms/sample - loss: 3.6785 - accuracy: 0.6821 - val_
Epoch 2/40
758/758 [=====] - 0s 286us/sample - loss: 2.3124 - accuracy: 0.7718 - va
Epoch 3/40
758/758 [=====] - 0s 283us/sample - loss: 1.2686 - accuracy: 0.7718 - va
Epoch 4/40
758/758 [=====] - 0s 327us/sample - loss: 1.1388 - accuracy: 0.7718 - va
Epoch 5/40
758/758 [=====] - 0s 348us/sample - loss: 1.0823 - accuracy: 0.7718 - va
Epoch 6/40
758/758 [=====] - 0s 298us/sample - loss: 1.0622 - accuracy: 0.7718 - va
Epoch 7/40
758/758 [=====] - 0s 300us/sample - loss: 1.0503 - accuracy: 0.7718 - va
Epoch 8/40
758/758 [=====] - 0s 308us/sample - loss: 1.0449 - accuracy: 0.7718 - va
Epoch 9/40
758/758 [=====] - 0s 280us/sample - loss: 1.0382 - accuracy: 0.7718 - va
Epoch 10/40
758/758 [=====] - 0s 286us/sample - loss: 1.0350 - accuracy: 0.7718 - va
Epoch 11/40
758/758 [=====] - 0s 291us/sample - loss: 1.0334 - accuracy: 0.7718 - va
Epoch 12/40
758/758 [=====] - 0s 322us/sample - loss: 1.0319 - accuracy: 0.7718 - va
Epoch 13/40
758/758 [=====] - 0s 323us/sample - loss: 1.0290 - accuracy: 0.7718 - va
Epoch 14/40
758/758 [=====] - 0s 288us/sample - loss: 1.0291 - accuracy: 0.7718 - va
Epoch 15/40
758/758 [=====] - 0s 274us/sample - loss: 1.0271 - accuracy: 0.7718 - va
```

Epoch 16/40
 758/758 [=====] - 0s 275us/sample - loss: 1.0234 - accuracy: 0.7718 - va
 Epoch 17/40
 758/758 [=====] - 0s 285us/sample - loss: 1.0145 - accuracy: 0.7718 - va
 Epoch 18/40
 758/758 [=====] - 0s 285us/sample - loss: 0.9961 - accuracy: 0.7718 - va
 Epoch 19/40
 758/758 [=====] - 0s 282us/sample - loss: 0.9568 - accuracy: 0.7718 - va
 Epoch 20/40
 758/758 [=====] - 0s 275us/sample - loss: 0.8873 - accuracy: 0.7718 - va
 Epoch 21/40
 758/758 [=====] - 0s 278us/sample - loss: 0.9047 - accuracy: 0.7718 - va
 Epoch 22/40
 758/758 [=====] - 0s 294us/sample - loss: 1.0311 - accuracy: 0.7718 - va
 Epoch 23/40
 758/758 [=====] - 0s 289us/sample - loss: 1.0250 - accuracy: 0.7718 - va
 Epoch 24/40
 758/758 [=====] - 0s 287us/sample - loss: 1.0251 - accuracy: 0.7718 - va
 Epoch 25/40
 758/758 [=====] - 0s 336us/sample - loss: 1.0201 - accuracy: 0.7718 - va
 Epoch 26/40
 758/758 [=====] - 0s 329us/sample - loss: 1.0066 - accuracy: 0.7718 - va
 Epoch 27/40
 758/758 [=====] - 0s 432us/sample - loss: 0.9466 - accuracy: 0.7718 - va
 Epoch 28/40
 758/758 [=====] - 0s 377us/sample - loss: 0.8134 - accuracy: 0.7810 - va
 Epoch 29/40
 758/758 [=====] - 0s 310us/sample - loss: 0.9883 - accuracy: 0.7836 - va
 Epoch 30/40
 758/758 [=====] - 0s 288us/sample - loss: 1.0242 - accuracy: 0.7718 - va
 Epoch 31/40
 758/758 [=====] - 0s 267us/sample - loss: 1.0188 - accuracy: 0.7718 - va
 Epoch 32/40
 758/758 [=====] - 0s 385us/sample - loss: 1.0178 - accuracy: 0.7718 - va
 Epoch 33/40
 758/758 [=====] - 0s 400us/sample - loss: 1.0166 - accuracy: 0.7718 - va
 Epoch 34/40
 758/758 [=====] - 0s 435us/sample - loss: 1.0162 - accuracy: 0.7718 - va
 Epoch 35/40
 758/758 [=====] - 0s 381us/sample - loss: 1.0164 - accuracy: 0.7718 - va
 Epoch 36/40
 758/758 [=====] - 0s 409us/sample - loss: 1.0155 - accuracy: 0.7718 - va
 Epoch 37/40
 758/758 [=====] - 0s 446us/sample - loss: 1.0130 - accuracy: 0.7718 - va
 Epoch 38/40
 758/758 [=====] - 0s 426us/sample - loss: 1.0148 - accuracy: 0.7718 - va
 Epoch 39/40
 758/758 [=====] - 0s 340us/sample - loss: 0.9521 - accuracy: 0.7718 - va
 Epoch 40/40

758/758 [=====] - 0s 261us/sample - loss: 0.7970 - accuracy: 0.8153 - va

```
In [24]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
```

```
print('\nTest accuracy:', test_acc)
```

272/1 - 0s - loss: 1.0245 - accuracy: 0.8824

Test accuracy: 0.88235295

- Modificando inicializador

```
In [25]: def deep_lstm():
        model = Sequential()
        model.add(Embedding(num_words, emb_dim))
        model.add(LSTM(32, return_sequences = True, recurrent_initializer='glorot_uniform'))
        model.add(LSTM(32, return_sequences = False, recurrent_initializer='glorot_uniform'))
        model.add(Dense(46, activation = 'softmax'))
        model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['ac

        return model
```

```
In [26]: model = deep_lstm()
        model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, None, 64)	512000
lstm_3 (LSTM)	(None, None, 32)	12416
lstm_4 (LSTM)	(None, 32)	8320
dense_3 (Dense)	(None, 46)	1518
Total params: 534,254		
Trainable params: 534,254		
Non-trainable params: 0		

```
In [27]: history2 = model.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = 40, verbose = verbose)
```

Train on 758 samples, validate on 326 samples

Epoch 1/40

758/758 [=====] - 4s 5ms/sample - loss: 3.5942 - accuracy: 0.6478 - val_

Epoch 2/40
758/758 [=====] - 0s 397us/sample - loss: 2.1352 - accuracy: 0.7718 - va
Epoch 3/40
758/758 [=====] - 0s 412us/sample - loss: 1.2587 - accuracy: 0.7718 - va
Epoch 4/40
758/758 [=====] - 0s 296us/sample - loss: 1.1126 - accuracy: 0.7718 - va
Epoch 5/40
758/758 [=====] - 0s 352us/sample - loss: 1.0777 - accuracy: 0.7718 - va
Epoch 6/40
758/758 [=====] - 0s 364us/sample - loss: 1.0595 - accuracy: 0.7718 - va
Epoch 7/40
758/758 [=====] - 0s 442us/sample - loss: 1.0511 - accuracy: 0.7718 - va
Epoch 8/40
758/758 [=====] - 0s 438us/sample - loss: 1.0439 - accuracy: 0.7718 - va
Epoch 9/40
758/758 [=====] - 0s 451us/sample - loss: 1.0396 - accuracy: 0.7718 - va
Epoch 10/40
758/758 [=====] - 0s 416us/sample - loss: 1.0362 - accuracy: 0.7718 - va
Epoch 11/40
758/758 [=====] - 0s 328us/sample - loss: 1.0344 - accuracy: 0.7718 - va
Epoch 12/40
758/758 [=====] - 0s 261us/sample - loss: 1.0335 - accuracy: 0.7718 - va
Epoch 13/40
758/758 [=====] - 0s 263us/sample - loss: 1.0302 - accuracy: 0.7718 - va
Epoch 14/40
758/758 [=====] - 0s 261us/sample - loss: 1.0296 - accuracy: 0.7718 - va
Epoch 15/40
758/758 [=====] - 0s 271us/sample - loss: 1.0282 - accuracy: 0.7718 - va
Epoch 16/40
758/758 [=====] - 0s 261us/sample - loss: 1.0273 - accuracy: 0.7718 - va
Epoch 17/40
758/758 [=====] - 0s 259us/sample - loss: 1.0265 - accuracy: 0.7718 - va
Epoch 18/40
758/758 [=====] - 0s 261us/sample - loss: 1.0272 - accuracy: 0.7718 - va
Epoch 19/40
758/758 [=====] - 0s 268us/sample - loss: 1.0248 - accuracy: 0.7718 - va
Epoch 20/40
758/758 [=====] - 0s 498us/sample - loss: 1.0261 - accuracy: 0.7718 - va
Epoch 21/40
758/758 [=====] - 0s 396us/sample - loss: 1.0247 - accuracy: 0.7718 - va
Epoch 22/40
758/758 [=====] - 0s 260us/sample - loss: 1.0237 - accuracy: 0.7718 - va
Epoch 23/40
758/758 [=====] - 0s 253us/sample - loss: 1.0231 - accuracy: 0.7718 - va
Epoch 24/40
758/758 [=====] - 0s 258us/sample - loss: 1.0235 - accuracy: 0.7718 - va
Epoch 25/40
758/758 [=====] - 0s 259us/sample - loss: 1.0217 - accuracy: 0.7718 - va
Epoch 26/40

```

758/758 [=====] - 0s 312us/sample - loss: 1.0144 - accuracy: 0.7718 - va
Epoch 27/40
758/758 [=====] - 0s 269us/sample - loss: 0.9809 - accuracy: 0.7718 - va
Epoch 28/40
758/758 [=====] - 0s 261us/sample - loss: 0.8520 - accuracy: 0.7770 - va
Epoch 29/40
758/758 [=====] - 0s 267us/sample - loss: 0.7102 - accuracy: 0.8602 - va
Epoch 30/40
758/758 [=====] - 0s 292us/sample - loss: 0.6853 - accuracy: 0.8417 - va
Epoch 31/40
758/758 [=====] - 0s 293us/sample - loss: 0.6408 - accuracy: 0.8562 - va
Epoch 32/40
758/758 [=====] - 0s 275us/sample - loss: 0.5880 - accuracy: 0.8641 - va
Epoch 33/40
758/758 [=====] - 0s 390us/sample - loss: 0.5779 - accuracy: 0.8615 - va
Epoch 34/40
758/758 [=====] - 0s 406us/sample - loss: 0.5824 - accuracy: 0.8654 - va
Epoch 35/40
758/758 [=====] - 0s 270us/sample - loss: 0.5104 - accuracy: 0.8852 - va
Epoch 36/40
758/758 [=====] - 0s 308us/sample - loss: 0.4854 - accuracy: 0.8945 - va
Epoch 37/40
758/758 [=====] - 0s 323us/sample - loss: 0.4664 - accuracy: 0.8945 - va
Epoch 38/40
758/758 [=====] - 0s 434us/sample - loss: 0.4511 - accuracy: 0.8905 - va
Epoch 39/40
758/758 [=====] - 0s 414us/sample - loss: 0.4427 - accuracy: 0.8892 - va
Epoch 40/40
758/758 [=====] - 0s 457us/sample - loss: 0.4296 - accuracy: 0.9037 - va

```

```
In [28]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
```

```
print('\nTest accuracy:', test_acc)
```

```
272/1 - 0s - loss: 0.9204 - accuracy: 0.8934
```

```
Test accuracy: 0.8933824
```

- Recurrent dropout

```
In [29]: def deep_lstm():
    model = Sequential()
    model.add(Embedding(num_words, emb_dim))
    model.add(LSTM(32, return_sequences = True, recurrent_initializer='glorot_uniform',
        recurrent_dropout=0.1))
    model.add(LSTM(32, return_sequences = False, recurrent_initializer='glorot_uniform',
        recurrent_dropout=0.1))
    model.add(Dense(46, activation = 'softmax'))
```

```

model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['ac

return model

```

```

In [30]: model = deep_lstm()
         model.summary()

```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(None, None, 64)	512000
lstm_5 (LSTM)	(None, None, 32)	12416
lstm_6 (LSTM)	(None, 32)	8320
dense_4 (Dense)	(None, 46)	1518
Total params: 534,254		
Trainable params: 534,254		
Non-trainable params: 0		

```

In [31]: history3 = model.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = epoch, verbose = verbose)

```

Train on 758 samples, validate on 326 samples

```

Epoch 1/40
758/758 [=====] - 4s 5ms/sample - loss: 3.4562 - accuracy: 0.6939 - val_
Epoch 2/40
758/758 [=====] - 1s 2ms/sample - loss: 1.8847 - accuracy: 0.7718 - val_
Epoch 3/40
758/758 [=====] - 1s 2ms/sample - loss: 1.2168 - accuracy: 0.7718 - val_
Epoch 4/40
758/758 [=====] - 1s 2ms/sample - loss: 1.0976 - accuracy: 0.7718 - val_
Epoch 5/40
758/758 [=====] - 1s 2ms/sample - loss: 1.0700 - accuracy: 0.7718 - val_
Epoch 6/40
758/758 [=====] - 1s 2ms/sample - loss: 1.0546 - accuracy: 0.7718 - val_
Epoch 7/40
758/758 [=====] - 1s 2ms/sample - loss: 1.0440 - accuracy: 0.7718 - val_
Epoch 8/40
758/758 [=====] - 1s 2ms/sample - loss: 1.0232 - accuracy: 0.7718 - val_
Epoch 9/40
758/758 [=====] - 1s 2ms/sample - loss: 0.9878 - accuracy: 0.7718 - val_
Epoch 10/40
758/758 [=====] - 2s 2ms/sample - loss: 0.9038 - accuracy: 0.7718 - val_
Epoch 11/40

```



```

758/758 [=====] - 1s 2ms/sample - loss: 0.8087 - accuracy: 0.7731 - val_
Epoch 12/40
758/758 [=====] - 1s 2ms/sample - loss: 0.7716 - accuracy: 0.8232 - val_
Epoch 13/40
758/758 [=====] - 1s 2ms/sample - loss: 0.7254 - accuracy: 0.8245 - val_
Epoch 14/40
758/758 [=====] - 1s 2ms/sample - loss: 0.6757 - accuracy: 0.8509 - val_
Epoch 15/40
758/758 [=====] - 1s 2ms/sample - loss: 0.6354 - accuracy: 0.8654 - val_
Epoch 16/40
758/758 [=====] - 1s 2ms/sample - loss: 0.6219 - accuracy: 0.8602 - val_
Epoch 17/40
758/758 [=====] - 1s 2ms/sample - loss: 0.6038 - accuracy: 0.8628 - val_
Epoch 18/40
758/758 [=====] - 1s 2ms/sample - loss: 0.5822 - accuracy: 0.8615 - val_
Epoch 19/40
758/758 [=====] - 1s 2ms/sample - loss: 0.5577 - accuracy: 0.8694 - val_
Epoch 20/40
758/758 [=====] - 1s 2ms/sample - loss: 0.5408 - accuracy: 0.8668 - val_
Epoch 21/40
758/758 [=====] - 1s 2ms/sample - loss: 0.5221 - accuracy: 0.8668 - val_
Epoch 22/40
758/758 [=====] - 1s 2ms/sample - loss: 0.5132 - accuracy: 0.8654 - val_
Epoch 23/40
758/758 [=====] - 1s 2ms/sample - loss: 0.5049 - accuracy: 0.8707 - val_
Epoch 24/40
758/758 [=====] - 1s 2ms/sample - loss: 0.4926 - accuracy: 0.8707 - val_
Epoch 25/40
758/758 [=====] - 1s 2ms/sample - loss: 0.5031 - accuracy: 0.8654 - val_
Epoch 26/40
758/758 [=====] - 1s 2ms/sample - loss: 0.4874 - accuracy: 0.8799 - val_
Epoch 27/40
758/758 [=====] - 1s 2ms/sample - loss: 0.4686 - accuracy: 0.8694 - val_
Epoch 28/40
758/758 [=====] - 1s 2ms/sample - loss: 0.4507 - accuracy: 0.8918 - val_
Epoch 29/40
758/758 [=====] - 1s 2ms/sample - loss: 0.4365 - accuracy: 0.8865 - val_
Epoch 30/40
758/758 [=====] - 1s 2ms/sample - loss: 0.4259 - accuracy: 0.8945 - val_
Epoch 31/40
758/758 [=====] - 1s 2ms/sample - loss: 0.4198 - accuracy: 0.8918 - val_
Epoch 32/40
758/758 [=====] - 1s 2ms/sample - loss: 0.4244 - accuracy: 0.8918 - val_
Epoch 33/40
758/758 [=====] - 1s 2ms/sample - loss: 0.4049 - accuracy: 0.9011 - val_
Epoch 34/40
758/758 [=====] - 1s 2ms/sample - loss: 0.3981 - accuracy: 0.9024 - val_
Epoch 35/40
758/758 [=====] - 1s 2ms/sample - loss: 0.3874 - accuracy: 0.9050 - val_

```

```

Epoch 36/40
758/758 [=====] - 1s 2ms/sample - loss: 0.3764 - accuracy: 0.9169 - val_
Epoch 37/40
758/758 [=====] - 1s 2ms/sample - loss: 0.3711 - accuracy: 0.9116 - val_
Epoch 38/40
758/758 [=====] - 1s 2ms/sample - loss: 0.3615 - accuracy: 0.9129 - val_
Epoch 39/40
758/758 [=====] - 1s 2ms/sample - loss: 0.3595 - accuracy: 0.9195 - val_
Epoch 40/40
758/758 [=====] - 1s 2ms/sample - loss: 0.3559 - accuracy: 0.9182 - val_

```

```
In [32]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
```

```
print('\nTest accuracy:', test_acc)
```

```
272/1 - 0s - loss: 0.9561 - accuracy: 0.9118
```

```
Test accuracy: 0.9117647
```

```
In [34]: #plot
```

```
plt.figure(figsize=(10,9))
```

```
plt.plot(history1.history['accuracy'])
```

```
plt.plot(history1.history['val_accuracy'])
```

```
plt.plot(history2.history['accuracy'])
```

```
plt.plot(history2.history['val_accuracy'])
```

```
plt.plot(history3.history['accuracy'])
```

```
plt.plot(history3.history['val_accuracy'])
```

```
plt.legend(['Train orthogonal', 'Test orthogonal',
```

```
           'Train gloriot_uniform', 'Test gloriot_uniform',
```


```
           'Train recurrent_dropout/gloriot_uniform', 'Test recurrent_dropout/gloriot_uni
```

```
plt.title('accuracy')
```

```
plt.ylabel('accuracy')
```

```
plt.xlabel('epoch')
```

```
plt.show()
```



- Mejorar test accuracy
- Investigar que son las GRU e implementarlas
- Experimentar con diferente número de capas y neuronas, mejorar los resultados
- Experimentar otros inicializadores y diferentes valores de dropout
- Probar con otro dataset