

Análisis de Datos y Aprendizaje Máquina con Tensorflow 2.0: Redes neuronales recurrentes

2019/09/30

Deep Bidirectional RNN

Objetivo: Comprender las RNN Bidireccionales y sus efectos en el aprendizaje del modelo.
Conocer el comportamiento de los optimizadores en modelos recurrentes.

```
In [1]: from tensorflow.keras.models import Model
        from tensorflow.keras.layers import LSTM, Embedding, Dense, SimpleRNN, Bidirectional
        from tensorflow.keras.datasets import imdb
        from tensorflow.keras.models import Sequential

        from tensorflow.keras.preprocessing.sequence import pad_sequences
        import matplotlib.pyplot as plt
```

Reseñas de películas de IMDB

- Conjunto de datos de 25,000 críticas de películas de IMDB, etiquetadas por sentimiento (positivo / negativo).

```
In [2]: # numero de palabras
        num_words = 4000

        max_len = 100

        #dimensión embedding
        emb_dim = 64

        (x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=num_words)

        x_train = pad_sequences(x_train, maxlen=max_len, padding = 'post')
        x_test = pad_sequences(x_test, maxlen=max_len, padding = 'post')

        print(x_train.shape)
        print(x_test.shape)
```

```

print(y_train.shape)
print(y_test.shape)

(25000, 100)
(25000, 100)
(25000,)
(25000,)

```

```

In [3]: epoch = 4
        verbose = 1
        batch = 30

```

Cada palabra de la review esta identificada por un número

```

In [4]: print('Reseña')
        print(x_train[0])
        print('Etiqueta')
        print(y_train[0])

```

```

Reseña
[1415  33   6  22  12 215  28  77  52   5  14 407  16  82
   2   8   4 107 117   2  15 256   4   2   7 3766   5 723
  36  71  43 530 476  26 400 317  46   7   4   2 1029  13
 104  88   4 381  15 297  98  32 2071  56  26 141   6 194
   2  18   4 226  22  21 134 476  26 480   5 144  30   2
  18  51  36  28 224  92  25 104   4 226  65  16  38 1334
  88  12  16 283   5  16   2 113 103  32  15  16   2  19
 178  32]
Etiqueta
1

```

Palabras de reseña

```

In [5]: wordDict = {y:x for x,y in imdb.get_word_index().items()}
        res = []
        for index in x_train[0]:
            res.append(wordDict.get(index - 3))
        print('Reseña: ',res,'Longitud reseña: ', len(res))

```

```

Reseña:  ['cry', 'at', 'a', 'film', 'it', 'must', 'have', 'been', 'good', 'and', 'this', 'definit

```

Deep RNN

- Se compara el modelo con y sin regularización

```

In [6]: def deep_lstm():
        model = Sequential()
        model.add(Embedding(num_words, emb_dim))

```

```

model.add(LSTM(64, return_sequences = True))
model.add(LSTM(64, return_sequences = False))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

return model

```

```

In [7]: model = deep_lstm()
        model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, None, 64)	256000
lstm (LSTM)	(None, None, 64)	33024
lstm_1 (LSTM)	(None, 64)	33024
dense (Dense)	(None, 1)	65

Total params: 322,113
 Trainable params: 322,113
 Non-trainable params: 0

```

In [8]: history = model.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = epoch, verbose = verbose)

```

Train on 17500 samples, validate on 7500 samples

Epoch 1/4

17500/17500 [=====] - 13s 729us/sample - loss: 0.4694 - accuracy: 0.7726

Epoch 2/4

17500/17500 [=====] - 8s 468us/sample - loss: 0.3294 - accuracy: 0.8625

Epoch 3/4

17500/17500 [=====] - 8s 473us/sample - loss: 0.2651 - accuracy: 0.8971

Epoch 4/4

17500/17500 [=====] - 8s 441us/sample - loss: 0.2223 - accuracy: 0.9161

```

In [9]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)

```

```

print('\nTest accuracy:', test_acc)


```

25000/1 - 4s - loss: 0.5023 - accuracy: 0.8368

Test accuracy: 0.8368

```
In [10]: plt.figure(figsize=(10,9))
```

```
plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('accuracy')
plt.legend(['train', 'test'])
plt.grid()
plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('loss')
plt.legend(['train', 'test'])
plt.grid()
```



Bidirectional RNN

- Bidirectional RNNs lee las secuencias en ambas direcciones

```
In [11]: def bidirectional_lstm():

        model = Sequential()
        model.add(Embedding(num_words, emb_dim))
        model.add(Bidirectional(LSTM(64, return_sequences = False)))
        model.add(Dense(1, activation='sigmoid'))
        model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

        return model
```

```
In [12]: model = bidirectional_lstm()
        model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, None, 64)	256000
bidirectional (Bidirectional)	(None, 128)	66048
dense_1 (Dense)	(None, 1)	129

Total params: 322,177

Trainable params: 322,177

Non-trainable params: 0

```
In [13]: history = model.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = epoch, verbose = verbose)
```

Train on 17500 samples, validate on 7500 samples

Epoch 1/4

17500/17500 [=====] - 11s 636us/sample - loss: 0.4474 - accuracy: 0.7817

Epoch 2/4

17500/17500 [=====] - 8s 459us/sample - loss: 0.2988 - accuracy: 0.8778

Epoch 3/4

17500/17500 [=====] - 8s 470us/sample - loss: 0.2477 - accuracy: 0.8997

Epoch 4/4

17500/17500 [=====] - 8s 468us/sample - loss: 0.1998 - accuracy: 0.9220

```
In [14]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
```


```
        print('\nTest accuracy:', test_acc)
```

25000/1 - 4s - loss: 0.6446 - accuracy: 0.8137

Test accuracy: 0.81368

```
In [15]: plt.figure(figsize=(10,9))
```

```
plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('accuracy')
plt.legend(['train', 'test'])
plt.grid()
plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('loss')
plt.legend(['train', 'test'])
plt.grid()
```



Deep Bidirectional RNN

- Observar el efecto de la profundidad de la red en el entrenamiento

```
In [16]: def deep_bidirectional_lstm():
        model = Sequential()
        model.add(Embedding(num_words, emb_dim))
        model.add(Bidirectional(LSTM(64, return_sequences = True)))
        model.add(Bidirectional(LSTM(64, return_sequences = False)))
        model.add(Dense(1, activation='sigmoid'))
        model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

        return model
```

```
In [17]: model = deep_bidirectional_lstm()
        model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, None, 64)	256000
bidirectional_1 (Bidirection	(None, None, 128)	66048
bidirectional_2 (Bidirection	(None, 128)	98816
dense_2 (Dense)	(None, 1)	129
Total params: 420,993		
Trainable params: 420,993		
Non-trainable params: 0		

```
In [18]: history1 = model.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = epoch, verbose = verbose)
```

Train on 17500 samples, validate on 7500 samples

Epoch 1/4

17500/17500 [=====] - 19s 1ms/sample - loss: 0.4360 - accuracy: 0.7903 -

Epoch 2/4

17500/17500 [=====] - 14s 798us/sample - loss: 0.2995 - accuracy: 0.8735 -

Epoch 3/4

17500/17500 [=====] - 14s 773us/sample - loss: 0.2407 - accuracy: 0.9041 -

Epoch 4/4

17500/17500 [=====] - 14s 784us/sample - loss: 0.1764 - accuracy: 0.9350 -

```
In [19]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
```


```
print('\nTest accuracy:', test_acc)
```

25000/1 - 7s - loss: 0.4300 - accuracy: 0.8369

Test accuracy: 0.83688

```
In [20]: plt.figure(figsize=(10,9))
```

```
plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('accuracy')
plt.legend(['train', 'test'])
plt.grid()
plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('loss')
plt.legend(['train', 'test'])
plt.grid()
```



Optimización en Bidirectional RNN con regularización

- Los optimizadores tienen comportamientos diferentes en las diferentes arquitecturas de las redes
- La regularización como dropout se agrega con 'recurrent_dropout'

```
In [21]: def deep_bidirectional_lstm():
        model = Sequential()
        model.add(Embedding(num_words, emb_dim))
        model.add(Bidirectional(LSTM(64, return_sequences = True,
                                     recurrent_initializer='glorot_uniform',
                                     recurrent_dropout=0.1)))
        model.add(Bidirectional(LSTM(64, return_sequences = False,
                                     recurrent_initializer='glorot_uniform',
                                     recurrent_dropout=0.1)))
        model.add(Dense(1, activation='sigmoid'))
        model.compile(loss='binary_crossentropy', optimizer='rmsprop', metrics=['accuracy'])

        return model
```

```
In [22]: model = deep_bidirectional_lstm()
        model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, None, 64)	256000
bidirectional_3 (Bidirection	(None, None, 128)	66048
bidirectional_4 (Bidirection	(None, 128)	98816
dense_3 (Dense)	(None, 1)	129

=====
Total params: 420,993
Trainable params: 420,993
Non-trainable params: 0
=====

```
In [23]: history2 = model.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = epoch, verbose = verbose)
```

Train on 17500 samples, validate on 7500 samples

Epoch 1/4

17500/17500 [=====] - 156s 9ms/sample - loss: 0.4587 - accuracy: 0.7790

Epoch 2/4

17500/17500 [=====] - 148s 8ms/sample - loss: 0.3404 - accuracy: 0.8582

Epoch 3/4

17500/17500 [=====] - 150s 9ms/sample - loss: 0.2987 - accuracy: 0.8794

Epoch 4/4

17500/17500 [=====] - 154s 9ms/sample - loss: 0.2658 - accuracy: 0.8933

```
In [24]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
```

```
print('\nTest accuracy:', test_acc)
```

25000/1 - 58s - loss: 0.4047 - accuracy: 0.8495

Test accuracy: 0.84952

- RMSProp con regularización supera a las redes anteriores con Adam
- Averiguar si quitando la regularización se llega a el mismo resultado

```
In [25]: def deep_bidirectional_lstm():
    model = Sequential()
    model.add(Embedding(num_words, emb_dim))
    model.add(Bidirectional(LSTM(64, return_sequences = True,
                                recurrent_initializer='glorot_uniform',
                                recurrent_dropout=0.1)))
    model.add(Bidirectional(LSTM(64, return_sequences = False,
                                recurrent_initializer='glorot_uniform',
                                recurrent_dropout=0.1)))
    model.add(Dense(1, activation='sigmoid'))
    model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

    return model
```

```
In [26]: model = deep_bidirectional_lstm()
    model.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(None, None, 64)	256000
bidirectional_5 (Bidirection	(None, None, 128)	66048
bidirectional_6 (Bidirection	(None, 128)	98816
dense_4 (Dense)	(None, 1)	129

=====
Total params: 420,993
Trainable params: 420,993
Non-trainable params: 0
=====

```
In [27]: history3 = model.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = epoch, verbose = verbose)
```

Train on 17500 samples, validate on 7500 samples

Epoch 1/4

17500/17500 [=====] - 154s 9ms/sample - loss: 0.4406 - accuracy: 0.7915

Epoch 2/4

17500/17500 [=====] - 148s 8ms/sample - loss: 0.3120 - accuracy: 0.8707

Epoch 3/4

17500/17500 [=====] - 146s 8ms/sample - loss: 0.2529 - accuracy: 0.8973

Epoch 4/4

17500/17500 [=====] - 147s 8ms/sample - loss: 0.2028 - accuracy: 0.9210

```
In [28]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
```

```
print('\nTest accuracy:', test_acc)
```

25000/1 - 58s - loss: 0.3618 - accuracy: 0.8309

Test accuracy: 0.83092

```
In [29]: # plot
plt.figure(figsize=(10,9))
plt.plot(history1.history['accuracy'])
plt.plot(history1.history['val_accuracy'])


plt.plot(history2.history['accuracy'])
plt.plot(history2.history['val_accuracy'])

plt.plot(history3.history['accuracy'])
plt.plot(history3.history['val_accuracy'])

plt.legend(['Train Adam simple', 'Test Adam simple',
            'Train RMSProp dropout 0.1', 'Test RMSProp dropout 0.1',
            'Train Adam dropout 0.1', 'Test Adam dropout 0.1',])

plt.title('accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')

plt.show()
```



- Mejorar el accuracy de la red
- Probar Deep-Bidirectional-RNN con otro dataset
- Probar con otros optimizadores y diferentes valores de 'recurrent_dropout' y 'recurrent_initializer'