

# Análisis de Datos y Aprendizaje Máquina con Tensorflow 2.0: Redes neuronales convolucionales


2019/09/30

## 1 Autoencoders

- Objetivo: comprender el funcionamiento de los autoencoders en el procesamiento de imágenes y aplicarlos para reconstruir imágenes
- Referencia: <https://blog.keras.io/building-autoencoders-in-keras.html>

Modelo que transforma una dimensión a una dimensión codificada para después reconstruir los datos. Usado para eliminar ruido y principio de los modelos ent-to-end ó [Seq-to-seq](#)

Input                      Encoded                      Output



Un autoencoder consiste en un 'encoder' que codifica la entrada a una dimensión indicada y un 'decoder' el cual genera la entrada a partir de la codificación del encoder.

- Tipos de datos a los que se aplican autoencoders:
  - Imágenes
  - Audio
  - Video

Un autoencoder aprende una distribución de los datos para después generarlos.

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras

from tensorflow.keras.layers import Dense, Input
from tensorflow.keras.models import Model


mnist = keras.datasets.mnist

In [2]: (x_train, y_train), (x_test, y_test) = mnist.load_data()

In [3]: x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255

In [4]: fig, ax = plt.subplots(nrows=2, ncols=5, sharex=True, sharey=True,)
ax = ax.flatten()
for i in range(10):
    img = x_train[y_train == i][0].reshape(28, 28)
    ax[i].imshow(img, cmap='Greys', interpolation='nearest')

ax[0].set_xticks([])
ax[0].set_yticks([])
plt.tight_layout()
plt.show()
```



## 1.1 Modelo de autoencoder

- Autoencoder con una sola capa oculta

- Se define la dimensión a codificar, en este caso d=10

```
In [5]: epoch = 5
        verbose = 1
        batch = 50
```

- Indicar dimensión de codificación

```
In [6]: d = 10
```

```
In [7]: encoder = keras.models.Sequential([
        keras.layers.Flatten(input_shape=(28, 28)),
        keras.layers.Dense(d, activation="relu"),
    ])

    decoder = keras.models.Sequential([
        keras.layers.Dense(28 * 28, activation="sigmoid", input_shape=[d]),

        keras.layers.Reshape((28, 28))
    ])

    model = keras.models.Sequential([encoder, decoder])
```

```
In [8]: model.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
sequential (Sequential)	(None, 10)	7850
sequential_1 (Sequential)	(None, 28, 28)	8624

Total params: 16,474  
 Trainable params: 16,474  
 Non-trainable params: 0

```
In [9]: model.compile(loss="binary_crossentropy", optimizer='adam', metrics=['accuracy'])
```

```
        model.fit(x_train, x_train, epochs=epoch)
```

Train on 60000 samples

Epoch 1/5

60000/60000 [=====] - 4s 72us/sample - loss: 0.2082 - accuracy: 0.7959

Epoch 2/5

60000/60000 [=====] - 4s 62us/sample - loss: 0.1646 - accuracy: 0.8029

Epoch 3/5

60000/60000 [=====] - 4s 63us/sample - loss: 0.1596 - accuracy: 0.8042

Epoch 4/5

```
60000/60000 [=====] - 4s 61us/sample - loss: 0.1572 - accuracy: 0.8048
Epoch 5/5
60000/60000 [=====] - 4s 61us/sample - loss: 0.1560 - accuracy: 0.8051
```

```
Out[9]: <tensorflow.python.keras.callbacks.History at 0x7fe9a7796310>
```

```
In [10]: decoded_imgs = model.predict(x_test[:5])
```

```
In [11]: n = 5
        for i in range(n):
            # instancias de prueba
            ax = plt.subplot(2, n, i+1)
            plt.imshow(x_test[i].reshape(28,28), cmap='Greys', interpolation='nearest')

            ax.get_xaxis().set_visible(False)
            ax.get_yaxis().set_visible(False)


            # reconstrucción
            ax = plt.subplot(2, n, i+n+1)
            plt.imshow(decoded_imgs[i].reshape(28,28), cmap='Greys', )

            ax.get_xaxis().set_visible(False)
            ax.get_yaxis().set_visible(False)
        plt.show()
```



## 1.2 Deep Autoencoder

- Los Autoencoders generalmente son profundos y pueden estar conformados por diferentes tipos de capas



```
In [12]: encoder = keras.models.Sequential([
        keras.layers.Flatten(input_shape=(28, 28)),
        keras.layers.Dense(50, activation="relu"),
        keras.layers.Dense(d, activation="relu"),
    ])

    decoder = keras.models.Sequential([
        keras.layers.Dense(50, activation="relu", input_shape=[d]),
        keras.layers.Dense(28 * 28, activation="sigmoid"),
        keras.layers.Reshape((28, 28))
    ])

    model = keras.models.Sequential([encoder, decoder])
```

```
In [13]: model.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
sequential_3 (Sequential)	(None, 10)	39760
sequential_4 (Sequential)	(None, 28, 28)	40534
Total params: 80,294		
Trainable params: 80,294		
Non-trainable params: 0		

```
In [14]: model.compile(loss="binary_crossentropy", optimizer='adam', metrics=['accuracy'])
```

```
model.fit(x_train, x_train, epochs=epoch)
```

Train on 60000 samples

Epoch 1/5

```

60000/60000 [=====] - 4s 72us/sample - loss: 0.1864 - accuracy: 0.7977
Epoch 2/5
60000/60000 [=====] - 4s 68us/sample - loss: 0.1445 - accuracy: 0.8057
Epoch 3/5
60000/60000 [=====] - 4s 70us/sample - loss: 0.1368 - accuracy: 0.8071
Epoch 4/5
60000/60000 [=====] - 4s 66us/sample - loss: 0.1338 - accuracy: 0.8076
Epoch 5/5
60000/60000 [=====] - 4s 66us/sample - loss: 0.1318 - accuracy: 0.8080

```

Out[14]: <tensorflow.python.keras.callbacks.History at 0x7fe94c113850>

In [15]: decoded\_imgs = model.predict(x\_test[:5])

```

In [16]: n = 5
        for i in range(n):
            # instancias de prueba
            ax = plt.subplot(2, n, i+1)
            plt.imshow(x_test[i].reshape(28,28), cmap='Greys', interpolation='nearest')

            ax.get_xaxis().set_visible(False)
            ax.get_yaxis().set_visible(False)

            # reconstrucción
            ax = plt.subplot(2, n, i+n+1)
            plt.imshow(decoded_imgs[i].reshape(28,28), cmap='Greys', interpolation='nearest')

            ax.get_xaxis().set_visible(False)
            ax.get_yaxis().set_visible(False)
        plt.show()

```



### 1.3 Autoencoder Convolucional

- Los Autoencoders Convolucionales trabajan con capas convolucionales y pooling. Para el decoder se puede usar 'Conv2DTranspose'
- Los autoencoders variacionales generan nuevos datos. Los modelos generativos pueden tener muchas aplicaciones

```
In [17]: from tensorflow.keras.layers import Input, LeakyReLU, Reshape, Dense, Conv2D, MaxPooling2D
         from tensorflow.keras.models import Model, Sequential
         import tensorflow as tf
```

```
In [18]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
x_train = x_train.reshape(x_train.shape[0], 28, 28, 1).astype('float32') / 255
x_test = x_test.reshape(x_test.shape[0], 28, 28, 1).astype('float32') / 255
```

```
In [19]: print(x_train.shape)
         print(x_test.shape)
         print(y_train.shape)
         print(y_test.shape)
```

```
(60000, 28, 28, 1)
(10000, 28, 28, 1)
(60000,)
(10000,)
```

```
In [20]: def conv_ae():
         # encoder
         inputs = Input(shape=(28, 28, 1))
         x = Conv2D(16, 3, activation='relu', padding='same')(inputs)

         x = Conv2D(32, 3, activation='relu', padding='same')(x)

         x = Conv2D(64, 3, activation='relu', padding='same')(x)
         x = MaxPooling2D(padding='same')(x)

         x = Conv2D(128, 3, activation='relu', padding='same')(x)
         x = MaxPooling2D(padding='same')(x)

         encoded = Conv2D(256, 3, activation='relu', padding='same')(x)

         # decoder
         x = Conv2DTranspose(256, 3, strides=(2, 2), activation='relu', padding='same')(encoded)
         x = Conv2DTranspose(128, 3, strides=(2, 2), activation='relu', padding='same')(x)
```

```

x = Conv2DTranspose(64, 3, strides=(1, 1), activation='relu', padding='same')(x)
x = Conv2DTranspose(32, 3, strides=(1, 1), activation='relu', padding='same')(x)
x = Conv2DTranspose(16, 3, strides=(1, 1), activation='relu', padding='same')(x)
decoded = Conv2DTranspose(1, 3, activation='sigmoid', padding='same')(x)

autoencoder = Model(inputs, decoded)
autoencoder.compile(optimizer='adam',
                    loss='binary_crossentropy')
return autoencoder

autoencoder = conv_ae()

In [21]: autoencoder.fit(x_train, x_train, epochs=1)

Train on 60000 samples
60000/60000 [=====] - 76s 1ms/sample - loss: 0.0734

Out[21]: <tensorflow.python.keras.callbacks.History at 0x7fe9243428d0>

In [22]: decoded_imgs = autoencoder.predict(x_test[:5])

In [23]: n = 5
        for i in range(n):
            # instancias de prueba
            ax = plt.subplot(2, n, i+1)
            plt.imshow(x_test[i].reshape(28,28), cmap='Greys', interpolation='nearest')

            ax.get_xaxis().set_visible(False)
            ax.get_yaxis().set_visible(False)

            # reconstrucción
            ax = plt.subplot(2, n, i+n+1)
            plt.imshow(decoded_imgs[i].reshape(28,28), cmap='Greys', interpolation='nearest')

            ax.get_xaxis().set_visible(False)
            ax.get_yaxis().set_visible(False)
        plt.show()

```





- La reconstrucción de los datos es mucho mas potente con capas de convoluciones.
- Probar con diferente número de capas, neuronas y funciones de costo, argumentar los resultados de la reconstrucción
- Experimentar con otro dataset