Análisis de Datos y Aprendizaje Máquina con Tensorflow 2.0: Redes neuronales convolucionales

2019/09/30

Autoencoders

- Objetivo: comprender el funcionamiento de los autoencoders en el procesamiento de imágenes y aplicarlos para reconstruir imágenes

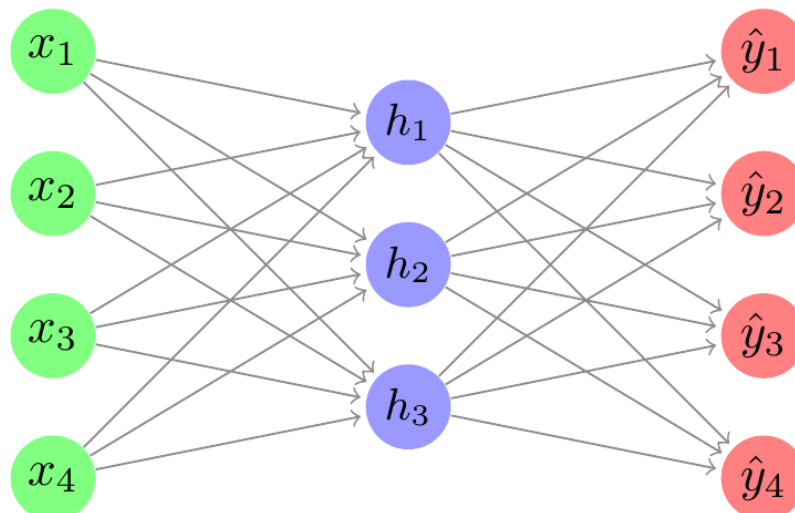https://blog.keras.io/building-autoencoders-in-keras.html

- Modelo que transforma una dimensión a una dimensión codificada para después reconstruir los datos
- Usada para eliminar ruido, principio de los modelos ent-to-end o seq-to-seq

Un autoencoder consiste en un 'encoder' que codifica la entrada a una dimensión indicada y un 'decoder' el cual genera la entrada a partir de la codificación del encoder.

- Tipos de datos a los que se aplican autoencoders:
  - Imagenes
  - Audio

– Video

Un autoencoder aprende una distribución de los datos para despúes generarlos.

```python
In [1]: import numpy as np
        import matplotlib.pyplot as plt
        import tensorflow as tf
        from tensorflow import keras

        from tensorflow.keras.layers import Dense, Input
        from tensorflow.keras.models import Model

        mnist = keras.datasets.mnist
```

```python
In [2]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```python
In [3]: x_train = x_train.astype('float32') / 255
        x_test = x_test.astype('float32') / 255
```

```python
In [4]: fig, ax = plt.subplots(nrows=2, ncols=5, sharex=True, sharey=True,)
        ax = ax.flatten()
        for i in range(10):
            img = x_train[y_train == i][0].reshape(28, 28)
            ax[i].imshow(img, cmap='Greys', interpolation='nearest')

        ax[0].set_xticks([])
        ax[0].set_yticks([])
        plt.tight_layout()
        plt.show()
```

Modelo de autoencoder

- Autoencoder con una sola capa oculta
- Se define la dimensión a codificar, en este caso d=10

```
In [5]: epoch = 5
        verbose = 1
        batch = 50
```

- Indicar dimensión de codificación

```
In [6]: d = 10
```

```
In [7]: encoder = keras.models.Sequential([
            keras.layers.Flatten(input_shape=(28, 28)),
            keras.layers.Dense(d, activation="relu"),
        ])

        decoder = keras.models.Sequential([
            keras.layers.Dense(28 * 28, activation="sigmoid", input_shape=[d]),

            keras.layers.Reshape((28, 28))
        ])

        model = keras.models.Sequential([encoder, decoder])
```

```
In [8]: model.summary()
```

```
Model: "sequential_2"

_____
Layer (type)                 Output Shape              Param #
=================================================================
sequential (Sequential)      (None, 10)                7850

_____
sequential_1 (Sequential)    (None, 28, 28)            8624
=================================================================
Total params: 16,474
Trainable params: 16,474
Non-trainable params: 0

_____
```

```
In [9]: model.compile(loss="binary_crossentropy", optimizer='adam', metrics=['accuracy'])

        model.fit(x_train, x_train, epochs=epoch)
```

```
Train on 60000 samples
Epoch 1/5
60000/60000 [==============================] - 4s 72us/sample - loss: 0.2082 - accuracy: 0.7959
Epoch 2/5
60000/60000 [==============================] - 4s 62us/sample - loss: 0.1646 - accuracy: 0.8029
```

```
Epoch 3/5
60000/60000 [==============================] - 4s 63us/sample - loss: 0.1596 - accuracy: 0.8042
Epoch 4/5
60000/60000 [==============================] - 4s 61us/sample - loss: 0.1572 - accuracy: 0.8048
Epoch 5/5
60000/60000 [==============================] - 4s 61us/sample - loss: 0.1560 - accuracy: 0.8051


Out[9]: <tensorflow.python.keras.callbacks.History at 0x7fe9a7796310>

In [10]: decoded_imgs = model.predict(x_test[:5])

In [11]: n = 5
         for i in range(n):
             # instancias de prueba
             ax = plt.subplot(2, n, i+1)
             plt.imshow(x_test[i].reshape(28,28), cmap='Greys', interpolation='nearest')


             ax.get_xaxis().set_visible(False)
             ax.get_yaxis().set_visible(False)

             # reconstrucción
             ax = plt.subplot(2, n, i+n+1)
             plt.imshow(decoded_imgs[i].reshape(28,28),cmap='Greys', )

             ax.get_xaxis().set_visible(False)
             ax.get_yaxis().set_visible(False)
         plt.show()
```
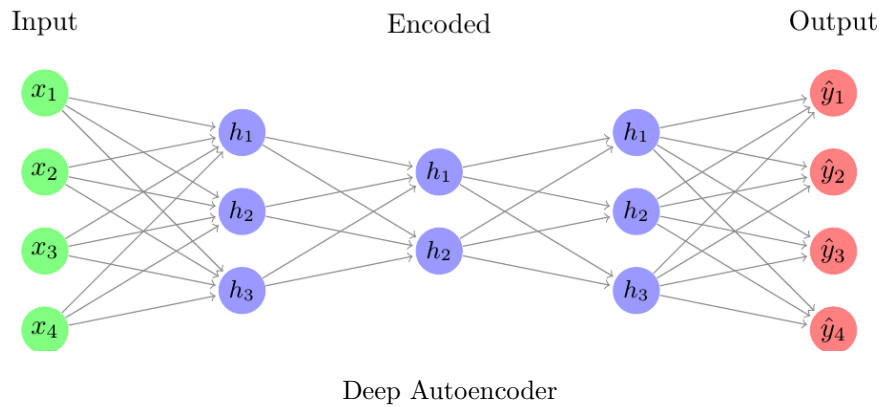
Input   Encoded   Output

Deep Autoencoder

Deep Autoencoder

- Los Autoencoders generalmente son profundos y pueden estar conformados por diferentes tipos de capas

```
In [12]: encoder = keras.models.Sequential([
             keras.layers.Flatten(input_shape=(28, 28)),
             keras.layers.Dense(50, activation="relu"),
             keras.layers.Dense(d, activation="relu"),
         ])

         decoder = keras.models.Sequential([
             keras.layers.Dense(50, activation="relu", input_shape=[d]),
             keras.layers.Dense(28 * 28, activation="sigmoid"),
             keras.layers.Reshape((28, 28))
         ])

         model = keras.models.Sequential([encoder, decoder])

In [13]: model.summary()

Model: "sequential_5"

_____
Layer (type)                 Output Shape              Param #
=================================================================
sequential_3 (Sequential)    (None, 10)                39760
_____
sequential_4 (Sequential)    (None, 28, 28)            40534
=================================================================
Total params: 80,294
Trainable params: 80,294
Non-trainable params: 0
_____
```

```
In [14]: model.compile(loss="binary_crossentropy", optimizer='adam', metrics=['accuracy'])

         model.fit(x_train, x_train, epochs=epoch)

Train on 60000 samples
Epoch 1/5
60000/60000 [==============================] - 4s 72us/sample - loss: 0.1864 - accuracy: 0.7977
Epoch 2/5
60000/60000 [==============================] - 4s 68us/sample - loss: 0.1445 - accuracy: 0.8057
Epoch 3/5
60000/60000 [==============================] - 4s 70us/sample - loss: 0.1368 - accuracy: 0.8071
Epoch 4/5
60000/60000 [==============================] - 4s 66us/sample - loss: 0.1338 - accuracy: 0.8076
Epoch 5/5
60000/60000 [==============================] - 4s 66us/sample - loss: 0.1318 - accuracy: 0.8080


Out[14]: <tensorflow.python.keras.callbacks.History at 0x7fe94c113850>

In [15]: decoded_imgs = model.predict(x_test[:5])

In [16]: n = 5
         for i in range(n):
             # instancias de prueba
             ax = plt.subplot(2, n, i+1)
             plt.imshow(x_test[i].reshape(28,28), cmap='Greys', interpolation='nearest')


             ax.get_xaxis().set_visible(False)
             ax.get_yaxis().set_visible(False)

             # reconstrucción
             ax = plt.subplot(2, n, i+n+1)
             plt.imshow(decoded_imgs[i].reshape(28,28),cmap='Greys', interpolation='nearest')

             ax.get_xaxis().set_visible(False)
             ax.get_yaxis().set_visible(False)
         plt.show()
```
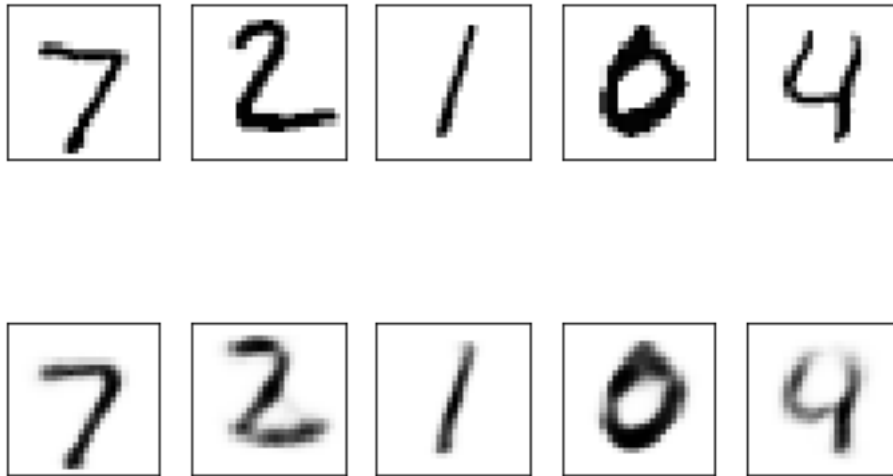
Autoencoder Convolucional

- Los Autoencoders Convolucionales trabajan con capas convolucionales y pooling. Para el decoder se puede usar 'Conv2DTranspose'
- Los autoencoders variacionales generan nuevos datos. Los modelos generativos pueden tener muchas aplicaciones

```
In [17]: from tensorflow.keras.layers import Input, LeakyReLU, Reshape, Dense, Conv2D, MaxPooling
         from tensorflow.keras.models import Model,Sequential
         import tensorflow as tf
```

```
In [18]: (x_train, y_train), (x_test, y_test) = mnist.load_data()


         x_train = x_train.reshape(x_train.shape[0], 28, 28, 1).astype('float32') / 255
         x_test = x_test.reshape(x_test.shape[0], 28, 28, 1).astype('float32') / 255
```

```
In [19]: print(x_train.shape)
         print(x_test.shape)
         print(y_train.shape)
         print(y_test.shape)
```

```
(60000, 28, 28, 1)
(10000, 28, 28, 1)
(60000,)
(10000,)
```

```
In [20]: def conv_ae():
             # encoder
```

```
            inputs = Input(shape=(28, 28, 1))
            x = Conv2D(16, 3, activation='relu', padding='same')(inputs)

            x = Conv2D(32, 3, activation='relu', padding='same')(x)

            x = Conv2D(64, 3, activation='relu', padding='same')(x)
            x = MaxPooling2D(padding='same')(x)

            x = Conv2D(128, 3, activation='relu', padding='same')(x)
            x = MaxPooling2D(padding='same')(x)

            encoded = Conv2D(256, 3, activation='relu', padding='same')(x)


            # decoder
            x = Conv2DTranspose(256, 3, strides=(2, 2), activation='relu', padding='same')(encod
            x = Conv2DTranspose(128, 3, strides=(2, 2), activation='relu', padding='same')(x)
            x = Conv2DTranspose(64, 3, strides=(1, 1), activation='relu', padding='same')(x)
            x = Conv2DTranspose(32, 3, strides=(1, 1), activation='relu', padding='same')(x)
            x = Conv2DTranspose(16, 3, strides=(1, 1), activation='relu', padding='same')(x)
            decoded = Conv2DTranspose(1, 3, activation='sigmoid', padding='same')(x)


            autoencoder = Model(inputs, decoded)
            autoencoder.compile(optimizer='adam',
                                loss='binary_crossentropy')
            return autoencoder


        autoencoder = conv_ae()

In [21]: autoencoder.fit(x_train, x_train, epochs=1)

Train on 60000 samples
60000/60000 [==============================] - 76s 1ms/sample - loss: 0.0734


Out[21]: <tensorflow.python.keras.callbacks.History at 0x7fe9243428d0>

In [22]: decoded_imgs = autoencoder.predict(x_test[:5])

In [23]: n = 5
        for i in range(n):
            # instancias de prueba
            ax = plt.subplot(2, n, i+1)
            plt.imshow(x_test[i].reshape(28,28), cmap='Greys', interpolation='nearest')


            ax.get_xaxis().set_visible(False)
            ax.get_yaxis().set_visible(False)
```

```python
# reconstrucción
ax = plt.subplot(2, n, i+n+1)
plt.imshow(decoded_imgs[i].reshape(28,28),cmap='Greys', interpolation='nearest')

ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()
```


- La reconstrucción de los datos es mucho mas potente con capas de convoluciones.

- Probar con diferente numero de capas, neuronas y funciones de costo, argumentar los resultados de la reconstrucción

- Experimentar con otro dataset