

Análisis de Datos y Aprendizaje Máquina con Tensorflow 2.0: Redes neuronales convolucionales

2019/09/30

Redes Neuronales Convolucionales (CNN)

- Objetivo: Conocer el tipo de capas de las CNN.
- Las redes convolucionales no solo se aplican a imágenes, también se pueden aplicar a caracteres o datos en el tiempo.

Convolución

- La convolución es una operación matemática generalmente denotada como $*$, en la que una función se aplica otra función, dando como resultado la combinación de las dos funciones.

Clasificar ropa

```
In [1]: import numpy as np
import matplotlib.pyplot as plt

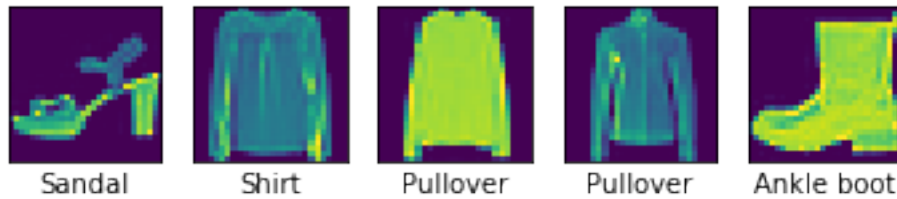
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import backend as K
K.clear_session()

fashion_mnist = keras.datasets.fashion_mnist

(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()

In [2]: class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress',
                        'Coat', 'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']

In [3]: for i in range(5):
    rand_image_idx = np.random.randint(0, y_train.shape[0])
    plt.subplot(1, 5, i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid('off')
    plt.imshow(x_train[rand_image_idx])
    plt.xlabel(class_names[y_train[rand_image_idx]])
plt.show()
```



```
In [4]: # escalar entre 0 y 1
x_train = x_train.reshape(x_train.shape[0], 28, 28, 1).astype('float32') / 255
x_test = x_test.reshape(x_test.shape[0], 28, 28, 1).astype('float32') / 255

print(x_train.shape) # (60000, 28, 28, 1)
print(x_test.shape)  # (10000, 28, 28, 1)

(60000, 28, 28, 1)
(10000, 28, 28, 1)
```

Obtener dimensiones

```
In [5]: # con 1: no se cuenta la primera dimensión
x, y, channel = x_train.shape[1:]

input_shape = (x, y, channel)
```

Crear modelo

```
In [6]: from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Dense, Activation, Flatten, Conv2D, MaxPooling2D

In [7]: epoch = 5
        verbose = 1
        batch = 50
```

Capa de convolución

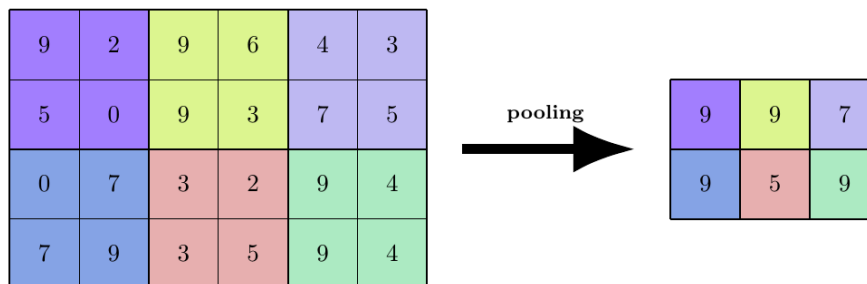
- En general se utiliza la convolución 2D para el procesamiento de imágenes
- En tamaño del filtro es igual al ancho y largo de los campos receptivos
- Pooling reduce el número de parámetros

```
In [8]: num_filters = 20
        filter_size = 3
        pool_size = 3
```

- Diferente notación para crear modelo

$$\begin{pmatrix} 0 & 1 & 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 \end{pmatrix} * \begin{pmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{pmatrix} = \begin{pmatrix} 0 & 4 & 2 & 5 \\ 1 & 2 & 4 & 3 \\ 1 & 2 & 3 & 4 \end{pmatrix}$$

Convolución



Pooling

```
In [9]: model = Sequential([
    Conv2D(num_filters, filter_size, input_shape=input_shape),
    MaxPooling2D(pool_size=pool_size),
    Flatten(),
    Dense(10, activation='softmax'),
])
```

Compiler

```
In [10]: model.compile('adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
In [11]: model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|------------------------------|--------------------|---------|
| conv2d (Conv2D) | (None, 26, 26, 20) | 200 |
| max_pooling2d (MaxPooling2D) | (None, 8, 8, 20) | 0 |
| flatten (Flatten) | (None, 1280) | 0 |
| dense (Dense) | (None, 10) | 12810 |


Total params: 13,010
 Trainable params: 13,010

Non-trainable params: 0

```
In [12]: from tensorflow.keras.utils import plot_model
```

```
In [13]: plot_model(model)
```

Out[13]:



Entrenamiento

```
In [14]: history = model.fit(x_train, y_train,
                             batch_size = batch,
                             validation_split=0.3,
                             epochs=epoch, verbose = verbose)
```

Train on 42000 samples, validate on 18000 samples

Epoch 1/5

42000/42000 [=====] - 5s 116us/sample - loss: 0.6021 - accuracy: 0.7926

Epoch 2/5

42000/42000 [=====] - 3s 60us/sample - loss: 0.4257 - accuracy: 0.8513

Epoch 3/5

42000/42000 [=====] - 3s 60us/sample - loss: 0.3909 - accuracy: 0.8632

Epoch 4/5

42000/42000 [=====] - 3s 60us/sample - loss: 0.3708 - accuracy: 0.8709

Epoch 5/5

42000/42000 [=====] - 2s 59us/sample - loss: 0.3559 - accuracy: 0.8758

```
In [15]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose = 0)
```

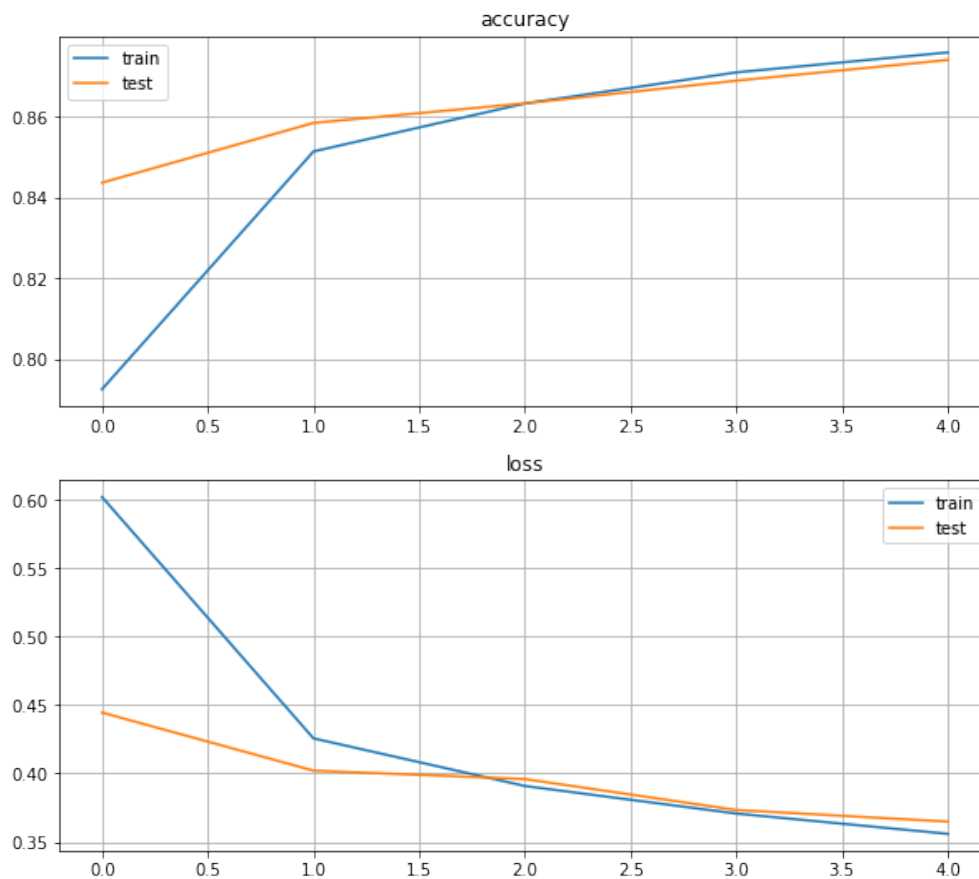
```
print('\nTest accuracy:', test_acc)
```

Test accuracy: 0.8669

```
In [16]: #plot
plt.figure(figsize=(10,9))

plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('accuracy')
plt.legend(['train', 'test'])
plt.grid()
plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('loss')
plt.legend(['train', 'test'])
plt.grid()

plt.show()
```



Diferente número de filtros

- Observar el número de parámetros y el efecto de los filtros y pooling en 'test accuracy'

```
In [17]: num_filters = 30
         filter_size = 3
         pool_size = 3

         model = Sequential([
             Conv2D(num_filters, filter_size, input_shape=input_shape),
             MaxPooling2D(pool_size=pool_size),
             Flatten(),
             Dense(10, activation='softmax'),
         ])

         model.compile('adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

         model.summary()
```

```

history = model.fit(x_train, y_train,
                    batch_size = batch,
                    validation_split=0.3,
                    epochs=epoch, verbose = verbose)

```

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|---------|
| conv2d_1 (Conv2D) | (None, 26, 26, 30) | 300 |
| max_pooling2d_1 (MaxPooling2D) | (None, 8, 8, 30) | 0 |
| flatten_1 (Flatten) | (None, 1920) | 0 |
| dense_1 (Dense) | (None, 10) | 19210 |

```

Total params: 19,510
Trainable params: 19,510
Non-trainable params: 0

```

Train on 42000 samples, validate on 18000 samples

```

Epoch 1/5
42000/42000 [=====] - 3s 70us/sample - loss: 0.5839 - accuracy: 0.7986 -
Epoch 2/5
42000/42000 [=====] - 3s 63us/sample - loss: 0.4235 - accuracy: 0.8513 -
Epoch 3/5
42000/42000 [=====] - 3s 64us/sample - loss: 0.3893 - accuracy: 0.8626 -
Epoch 4/5
42000/42000 [=====] - 3s 65us/sample - loss: 0.3639 - accuracy: 0.8720 -
Epoch 5/5
42000/42000 [=====] - 3s 64us/sample - loss: 0.3457 - accuracy: 0.8786 -

```

```

In [18]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose = 0)

```

```

print('\nTest accuracy:', test_acc)

```

Test accuracy: 0.8645

```

In [19]: #plot
plt.figure(figsize=(10,9))

plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('accuracy')

```

```

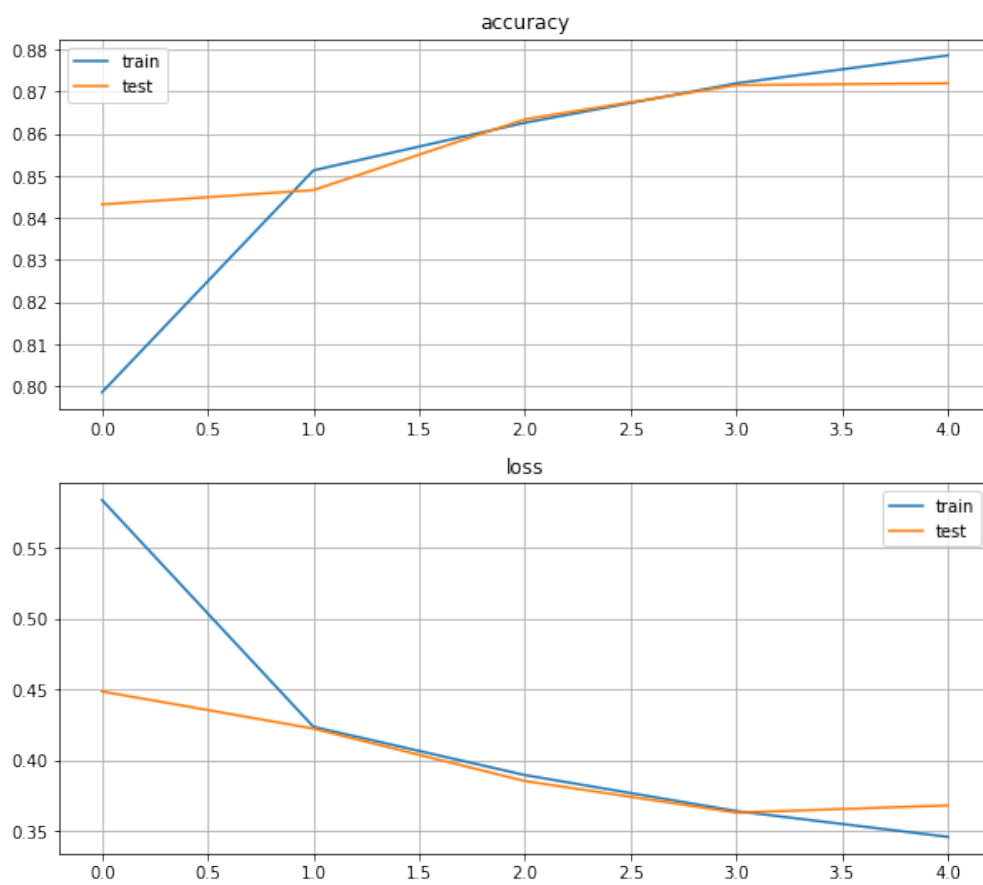
plt.legend(['train', 'test'])
plt.grid()
plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('loss')
plt.legend(['train', 'test'])
plt.grid()

```

```

plt.show()

```



Diferente tamaño de pooling

```

In [20]: num_filters = 30
        filter_size = 3

```



```

pool_size = 2

model = Sequential([
    Conv2D(num_filters, filter_size, input_shape=input_shape),
    MaxPooling2D(pool_size=pool_size),
    Flatten(),
    Dense(10, activation='softmax'),
])

model.compile('adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

model.summary()

history = model.fit(
    x_train, y_train,
    batch_size = batch,
    validation_split=0.3,
    epochs=epoch, verbose = verbose)

```

Model: "sequential_2"

| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|---------|
| conv2d_2 (Conv2D) | (None, 26, 26, 30) | 300 |
| max_pooling2d_2 (MaxPooling2D) | (None, 13, 13, 30) | 0 |
| flatten_2 (Flatten) | (None, 5070) | 0 |
| dense_2 (Dense) | (None, 10) | 50710 |

```

Total params: 51,010
Trainable params: 51,010
Non-trainable params: 0

```

Train on 42000 samples, validate on 18000 samples

```

Epoch 1/5
42000/42000 [=====] - 3s 75us/sample - loss: 0.5164 - accuracy: 0.8217 -
Epoch 2/5
42000/42000 [=====] - 3s 71us/sample - loss: 0.3831 - accuracy: 0.8673 -
Epoch 3/5
42000/42000 [=====] - 3s 70us/sample - loss: 0.3460 - accuracy: 0.8799 -
Epoch 4/5
42000/42000 [=====] - 3s 71us/sample - loss: 0.3211 - accuracy: 0.8883 -
Epoch 5/5
42000/42000 [=====] - 3s 73us/sample - loss: 0.3025 - accuracy: 0.8947 -

```

```
In [21]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose = 0)
```

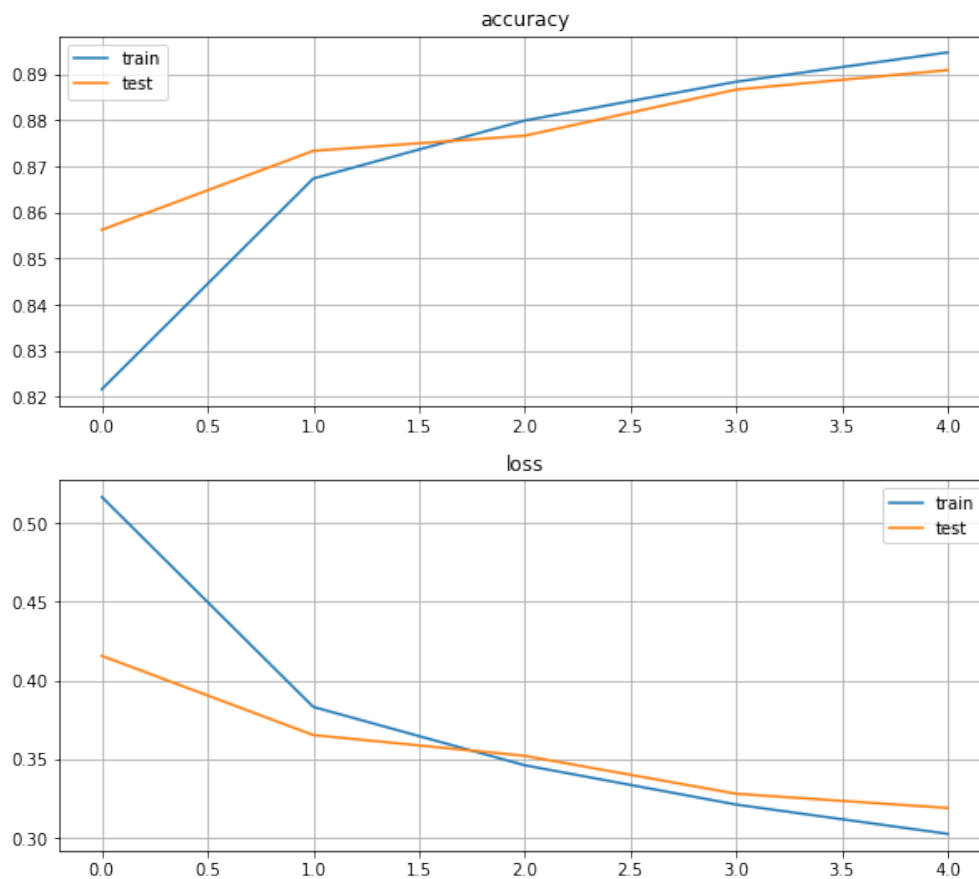
```
print('\nTest accuracy:', test_acc)
```

Test accuracy: 0.8782

```
In [22]: #plot
plt.figure(figsize=(10,9))

plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('accuracy')
plt.legend(['train', 'test'])
plt.grid()
plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('loss')
plt.legend(['train', 'test'])
plt.grid()

plt.show()
```



Probar predicciones del modelo

- 'argmax' retorna el elemento de mayor valor

```
In [23]: # Primeras 5 imagenes de test
         predictions = model.predict(x_test[:5])
```

```
         print(np.argmax(predictions, axis=1))
         p = np.argmax(predictions, axis=1)
```

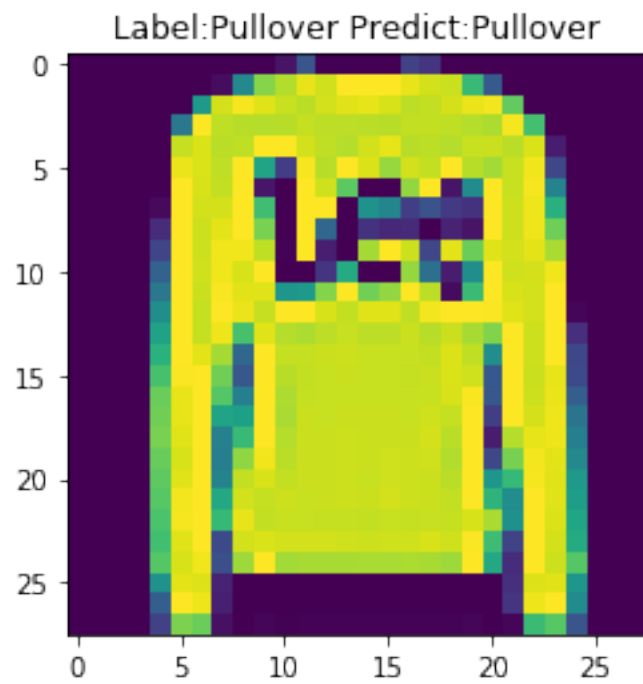
```
         print(y_test[:5])
```

```
[9 2 1 1 6]
```

```
[9 2 1 1 6]
```

```
In [24]: plt.imshow(np.squeeze(x_test[1]))
         plt.title('Label:' + class_names[int(y_test[1])])
```

```
+ ' Predict:' + class_names[int(p[1])])  
plt.show()
```



- Mejorar la arquitectura
- Probar con otro dataset