

# Análisis de Datos y Aprendizaje Máquina con Tensorflow 2.0: Redes neuronales convolucionales

2019/09/30

## 1 Redes Neuronales Convolucionales (CNN)

- Obetivo: Conocer el tipo de capas de las CNN.
- Las redes convolucionales no solo se aplican a imágenes, también se pueden aplicar a caracteres o datos en el tiempo.

### 1.1 Convolución

- La convolución es una operación matemática generalmente denotada como  $*$ , en la que una función se aplica otra función, dando como resultado la combinación de las dos funciones.

### 1.2 Clasificar ropa

```
In [1]: import numpy as np
import matplotlib.pyplot as plt


import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import backend as K
K.clear_session()

fashion_mnist = keras.datasets.fashion_mnist

(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()

In [2]: class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress',
                        'Coat', 'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']

In [3]: for i in range(5):
    rand_image_idx = np.random.randint(0, y_train.shape[0])
    plt.subplot(1, 5, i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid('off')
    plt.imshow(x_train[rand_image_idx])
    plt.xlabel(class_names[y_train[rand_image_idx]])
plt.show()
```



```
In [4]: # escalar entre 0 y 1
x_train = x_train.reshape(x_train.shape[0], 28, 28, 1).astype('float32') / 255
x_test = x_test.reshape(x_test.shape[0], 28, 28, 1).astype('float32') / 255

print(x_train.shape) # (60000, 28, 28, 1)
print(x_test.shape)  # (10000, 28, 28, 1)
```

(60000, 28, 28, 1)  
(10000, 28, 28, 1)

### 1.3 Obtener dimensiones

```
In [5]: # con 1: no se cuenta la primera dimensión
x, y, channel = x_train.shape[1:]

input_shape = (x, y, channel)
```

### 1.4 Crear modelo

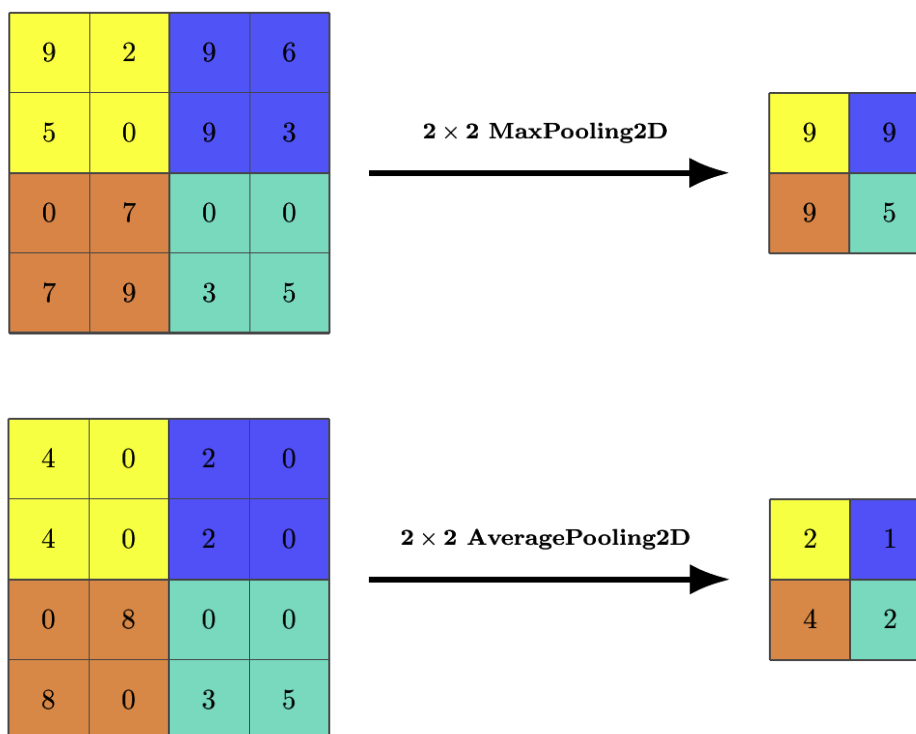
```
In [6]: from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Dense, Activation, Flatten, Conv2D, MaxPooling2D

In [7]: epoch = 5
        verbose = 1
        batch = 50
```

### 1.5 Capa de convolución

$$\begin{pmatrix} 0 & 1 & 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 \end{pmatrix} * \begin{pmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{pmatrix} = \begin{pmatrix} 0 & 4 & 2 & 5 \\ 1 & 2 & 4 & 3 \\ 1 & 2 & 3 & 4 \end{pmatrix}$$

- En general se utiliza la convolución 2D para el procesamiento de imágenes
- En tamaño del filtro es igual al ancho y largo de los campos receptivos
- Pooling reduce el número de parámetros



Imágenes generadas con <https://github.com/MartinThoma/LaTeX-examples>

```
In [8]: num_filters = 20
       filter_size = 3
       pool_size = 3
```

- Diferente notación para crear modelo

```
In [9]: model = Sequential([
        Conv2D(num_filters, filter_size, input_shape=input_shape),
        MaxPooling2D(pool_size=pool_size),
        Flatten(),
        Dense(10, activation='softmax'),
    ])
```

## 1.6 Compilar

```
In [10]: model.compile('adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
In [11]: model.summary()
```

Model: "sequential"

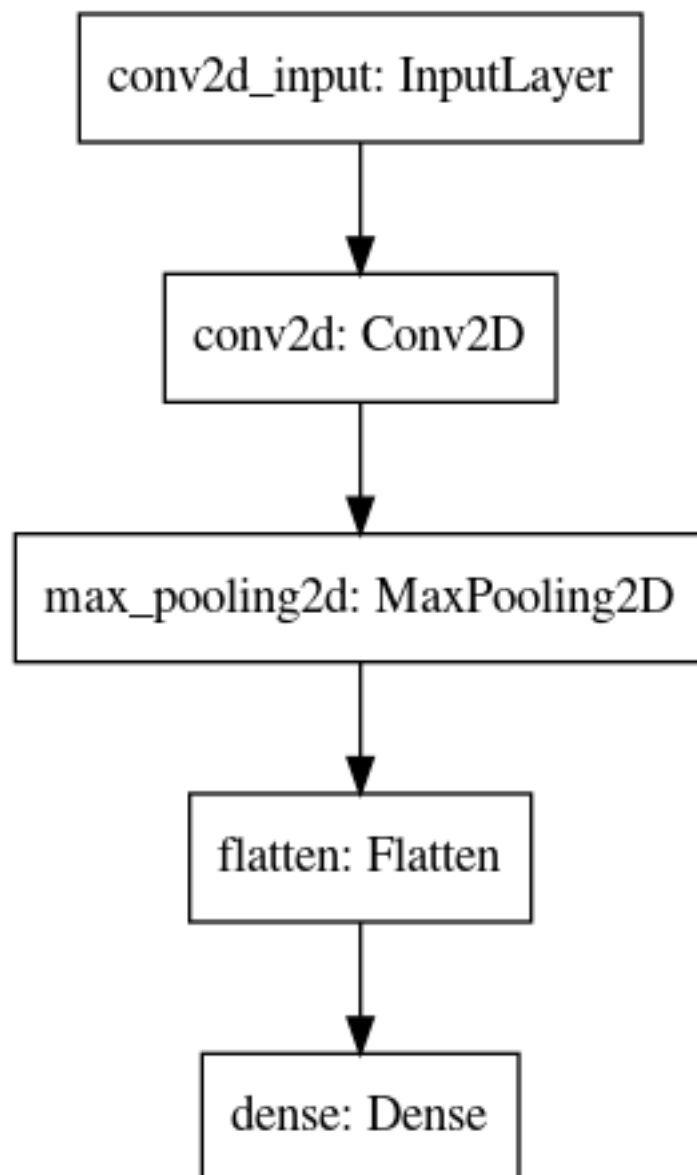
Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 20)	200
max_pooling2d (MaxPooling2D)	(None, 8, 8, 20)	0
flatten (Flatten)	(None, 1280)	0
dense (Dense)	(None, 10)	12810

Total params: 13,010  
Trainable params: 13,010  
Non-trainable params: 0

In [12]: `from tensorflow.keras.utils import plot_model`

In [13]: `plot_model(model)`

Out[13]:



## 1.7 Entrenamiento

```
In [14]: history = model.fit(x_train, y_train,  
                             batch_size = batch,  
                             validation_split=0.3,  
                             epochs=epoch, verbose = verbose)
```

Epoch 1/5

840/840 [=====] - 2s 3ms/step - loss: 0.6160 - accuracy: 0.7857 - val\_loss: 0.6160

```
Epoch 2/5
840/840 [=====] - 2s 2ms/step - loss: 0.4422 - accuracy: 0.8459 - val_lo
Epoch 3/5
840/840 [=====] - 2s 2ms/step - loss: 0.4054 - accuracy: 0.8584 - val_lo
Epoch 4/5
840/840 [=====] - 2s 2ms/step - loss: 0.3828 - accuracy: 0.8678 - val_lo
Epoch 5/5
840/840 [=====] - 2s 3ms/step - loss: 0.3634 - accuracy: 0.8725 - val_lo
```

```
In [15]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose = 0)


        print('\nTest accuracy:', test_acc)
```

Test accuracy: 0.8626999855041504

```
In [16]: #plot
plt.figure(figsize=(10,9))

plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('accuracy')
plt.legend(['train', 'test'])
plt.grid()
plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('loss')
plt.legend(['train', 'test'])
plt.grid()

plt.show()
```



## 1.8 Diferente número de filtros

- Observar el número de parámetros y el efecto de los filtros y pooling en 'test accuracy'

```
In [17]: num_filters = 30
         filter_size = 3
         pool_size = 3

         model = Sequential([
             Conv2D(num_filters, filter_size, input_shape=input_shape),
             MaxPooling2D(pool_size=pool_size),
             Flatten(),
             Dense(10, activation='softmax'),
         ])

         model.compile('adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

         model.summary()
```

```

history = model.fit(x_train, y_train,
                    batch_size = batch,
                    validation_split=0.3,
                    epochs=epoch, verbose = verbose)

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 30)	300
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 30)	0
flatten_1 (Flatten)	(None, 1920)	0
dense_1 (Dense)	(None, 10)	19210

Total params: 19,510

Trainable params: 19,510

Non-trainable params: 0

Epoch 1/5

840/840 [=====] - 2s 3ms/step - loss: 0.5705 - accuracy: 0.8058 - val\_loss: 0.4097

Epoch 2/5

840/840 [=====] - 2s 3ms/step - loss: 0.4097 - accuracy: 0.8581 - val\_loss: 0.3766

Epoch 3/5

840/840 [=====] - 2s 3ms/step - loss: 0.3766 - accuracy: 0.8692 - val\_loss: 0.3554

Epoch 4/5

840/840 [=====] - 2s 3ms/step - loss: 0.3554 - accuracy: 0.8757 - val\_loss: 0.3412

Epoch 5/5

840/840 [=====] - 2s 3ms/step - loss: 0.3412 - accuracy: 0.8812 - val\_loss: 0.3412

```

In [18]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose = 0)

```

```

print('\nTest accuracy:', test_acc)

```

Test accuracy: 0.871399998664856

```

In [19]: #plot

```

```

plt.figure(figsize=(10,9))

```

```

plt.subplot(211)

```

```

plt.plot(history.history['accuracy'])

```

```

plt.plot(history.history['val_accuracy'])

```

```

plt.title('accuracy')

```

```


plt.legend(['train', 'test'])

```



```
plt.grid()
plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('loss')
plt.legend(['train', 'test'])
plt.grid()
```

```
plt.show()
```



## 1.9 Strides y pooling

- Strides indica la unidad de desplazamiento del filtro. Si el valor es mayor a uno, los mapas de salida cuentan con un menor número de parámetros.

```

In [21]: num_filters = 30
         filter_size = 3

model = Sequential([
    Conv2D(num_filters, filter_size, input_shape=input_shape, strides=2),
    Flatten(),
    Dense(10, activation='softmax'),
])

model.compile('adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

model.summary()

history = model.fit(
    x_train, y_train,
    batch_size = batch,
    validation_split=0.3,
    epochs=epoch, verbose = verbose)

```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 13, 13, 30)	300
flatten_2 (Flatten)	(None, 5070)	0
dense_2 (Dense)	(None, 10)	50710

Total params: 51,010  
 Trainable params: 51,010  
 Non-trainable params: 0

```

-----
Epoch 1/5
840/840 [=====] - 2s 2ms/step - loss: 0.5653 - accuracy: 0.8016 - val_loss: 0.4540
Epoch 2/5
840/840 [=====] - 2s 2ms/step - loss: 0.4540 - accuracy: 0.8423 - val_loss: 0.4332
Epoch 3/5
840/840 [=====] - 2s 2ms/step - loss: 0.4332 - accuracy: 0.8490 - val_loss: 0.4227
Epoch 4/5
840/840 [=====] - 2s 2ms/step - loss: 0.4227 - accuracy: 0.8535 - val_loss: 0.4123
Epoch 5/5
840/840 [=====] - 2s 2ms/step - loss: 0.4123 - accuracy: 0.8555 - val_loss: 0.4123

```

```

In [22]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose = 0)

         print('\nTest accuracy:', test_acc)

```

Test accuracy: 0.8367000222206116

## 1.10 Diferente tamaño de pooling

- Se obtiene el mismo número de parámetros con strides=2 y sin capa pooling

```
In [23]: num_filters = 30
         filter_size = 3
         pool_size = 2

         model = Sequential([
             Conv2D(num_filters, filter_size, input_shape=input_shape),
             MaxPooling2D(pool_size=pool_size),
             Flatten(),
             Dense(10, activation='softmax'),
         ])

         model.compile('adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

         model.summary()

         history = model.fit(
             x_train, y_train,
             batch_size = batch,
             validation_split=0.3,
             epochs=epoch, verbose = verbose)
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 26, 26, 30)	300
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 30)	0
flatten_3 (Flatten)	(None, 5070)	0
dense_3 (Dense)	(None, 10)	50710

Total params: 51,010

Trainable params: 51,010

Non-trainable params: 0

Epoch 1/5

840/840 [=====] - 3s 3ms/step - loss: 0.5130 - accuracy: 0.8250 - val\_loss: 0.3769

Epoch 2/5

840/840 [=====] - 2s 3ms/step - loss: 0.3769 - accuracy: 0.8702 - val\_loss: 0.3396

Epoch 3/5

840/840 [=====] - 3s 3ms/step - loss: 0.3396 - accuracy: 0.8828 - val\_loss: 0.3396

```
Epoch 4/5
840/840 [=====] - 2s 3ms/step - loss: 0.3147 - accuracy: 0.8902 - val_lo
Epoch 5/5
840/840 [=====] - 2s 3ms/step - loss: 0.2952 - accuracy: 0.8976 - val_lo
```

```
In [24]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose = 0)

        print('\nTest accuracy:', test_acc)
```

```
Test accuracy: 0.8824999928474426
```

## 2 Probar predicciones del modelo

- ‘argmax’ retorna el elemento de mayor valor

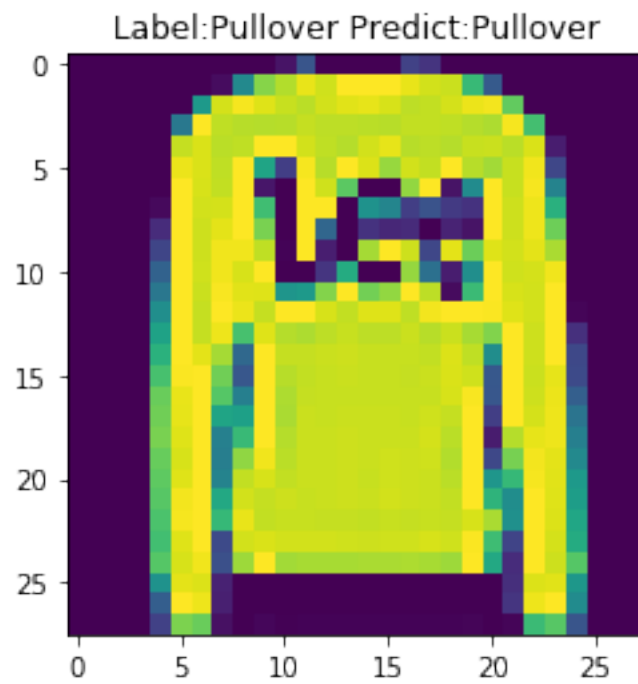
```
In [25]: # Primeras 5 imagenes de test
        predictions = model.predict(x_test[:5])

        print(np.argmax(predictions, axis=1))
        p = np.argmax(predictions, axis=1)

        print(y_test[:5])
```

```
[9 2 1 1 6]
[9 2 1 1 6]
```

```
In [26]: plt.imshow(np.squeeze(x_test[1]))
        plt.title('Label:' + class_names[int(y_test[1])]
                  + ' Predict:' + class_names[int(p[1])])
        plt.show()
```



- Mejorar la arquitectura
- Probar con otro dataset