

Análisis de Datos y Aprendizaje Máquina con Tensorflow 2.0: Redes neuronales convolucionales

2019/09/30

Computer Vision/Clasificación de objetos

<https://www.cs.toronto.edu/~kriz/cifar.html>

- Objetivo: Programar una red neuronal para reconocer objetos
- CIFAR es un acrónimo que significa Instituto Canadiense de Investigación Avanzada. El conjunto de datos CIFAR-10 consiste en 60000 imágenes a color de 32x32 para 10 clases de elementos.


```
In [9]: import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.datasets import cifar10

# leer dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

print('Train shape', x_train.shape, 'Target shape:', y_train.shape)
print('Test shape', x_test.shape, 'Target shape:', y_test.shape)
# plotea imagenes
for i in range(9):
    # define subplot
    plt.subplot(330 + 1 + i)
    plt.axis('off')
    plt.imshow(x_train[i])

plt.show()
```

```
Train shape (50000, 32, 32, 3) Target shape: (50000, 1)
Test shape (10000, 32, 32, 3) Target shape: (10000, 1)
```



Normalizar datos

- Número de clases es el número de neuronas que tendrá la capa de salida con 'softmax'

```
In [10]: import numpy as np
         clases = len(np.unique(y_train))
         clases
```

```
Out[10]: 10
```

```
In [11]: from tensorflow.keras.models import Sequential
         from tensorflow.keras.layers import Dense, Activation, Flatten, Conv2D, MaxPooling2D
         from tensorflow.keras.layers import Dropout, BatchNormalization
         # con 1: no se cuenta la primera dimensión
         x, y, channel = x_train.shape[1:]

         input_shape = (x, y, channel)

         # escalar entre 0 y 1
         # escalar entre 0 y 1
         x_train = x_train.reshape(x_train.shape[0], 32, 32, 3).astype('float32') / 255
         x_test = x_test.reshape(x_test.shape[0], 32, 32, 3).astype('float32') / 255

         print(x_train.shape) # (50000, 32, 32, 3)
         print(x_test.shape)  # (10000, 32, 32, 3)
```

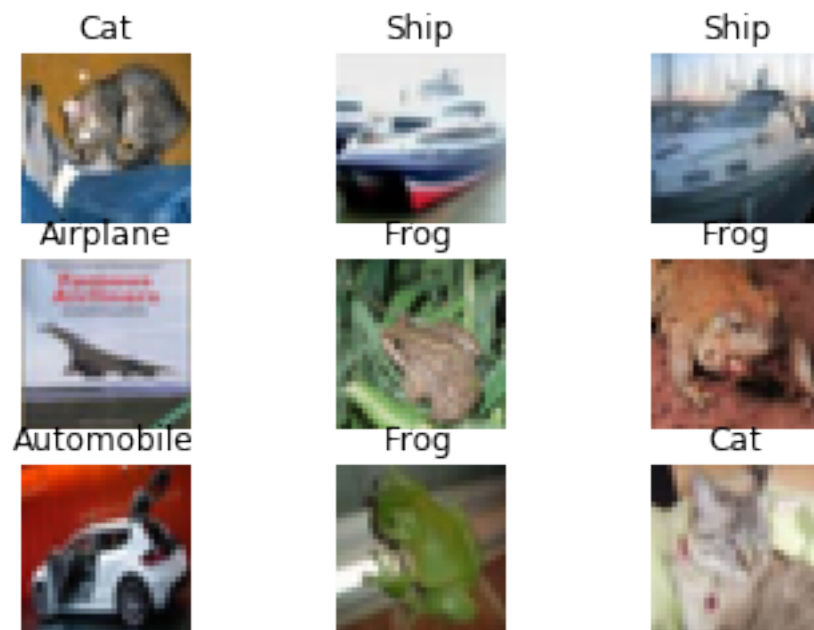
```
(50000, 32, 32, 3)
(10000, 32, 32, 3)
```

```
In [12]: int(y_test[i])
```

```
Out[12]: 3
```

```
In [13]: dict = {0:'Airplane', 1:'Automobile', 2:'Bird',
                 3:'Cat', 4:'Deer', 5:'Dog', 6:'Frog', 7:'Horse',
                 8:'Ship', 9:'Truck'}
for i in range(9):
    # define subplot
    plt.subplot(330 + 1 + i)
    plt.imshow(x_test[i])
    plt.axis('off')
    plt.title( dict[ int(y_test[i]) ] )
```

```
plt.show()
```



```
In [14]: epoch = 35
         verbose = 1
         batch = 50
```

Crear Modelo

```
In [15]: def deep_cnn():
    model = Sequential()

    model.add(Conv2D(20, (2, 2), padding='same', activation='relu',
                     input_shape=input_shape))
    model.add(Conv2D(20, (2, 2), activation='relu'))
    model.add(BatchNormalization(momentum=0.5))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Dropout(0.3))

    model.add(Conv2D(60, (2, 2), padding='same', activation='relu'))
    model.add(Conv2D(60, (2, 2), activation='relu'))
    model.add(BatchNormalization(momentum=0.5))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Dropout(0.3))

    model.add(Conv2D(120, (2, 2), padding='same', activation='relu'))
    model.add(Conv2D(120, (2, 2), activation='relu'))
    model.add(BatchNormalization(momentum=0.5))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Dropout(0.3))

    model.add(Flatten())
    model.add(Dense(64, activation='relu'))
    model.add(BatchNormalization(momentum=0.5))
    model.add(Dense(10, activation='softmax'))

    model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['acc'])

    return model
```

Entrenar modelo

- Nota: El entrenamiento necesita mucho poder de cómputo

```
In [16]: model1 = deep_cnn()
        model1.summary()

        history = model1.fit(x_train, y_train, batch_size = batch, validation_split = 0.3,
                             epochs = epoch, verbose = verbose)
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 32, 32, 20)	260
conv2d_7 (Conv2D)	(None, 31, 31, 20)	1620

batch_normalization_4 (Batch Normalization)	(None, 31, 31, 20)	80
max_pooling2d_3 (MaxPooling2D)	(None, 15, 15, 20)	0
dropout_3 (Dropout)	(None, 15, 15, 20)	0
conv2d_8 (Conv2D)	(None, 15, 15, 60)	4860
conv2d_9 (Conv2D)	(None, 14, 14, 60)	14460
batch_normalization_5 (Batch Normalization)	(None, 14, 14, 60)	240
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 60)	0
dropout_4 (Dropout)	(None, 7, 7, 60)	0
conv2d_10 (Conv2D)	(None, 7, 7, 120)	28920
conv2d_11 (Conv2D)	(None, 6, 6, 120)	57720
batch_normalization_6 (Batch Normalization)	(None, 6, 6, 120)	480
max_pooling2d_5 (MaxPooling2D)	(None, 3, 3, 120)	0
dropout_5 (Dropout)	(None, 3, 3, 120)	0
flatten_1 (Flatten)	(None, 1080)	0
dense_2 (Dense)	(None, 64)	69184
batch_normalization_7 (Batch Normalization)	(None, 64)	256
dense_3 (Dense)	(None, 10)	650

=====
Total params: 178,730

Trainable params: 178,202

Non-trainable params: 528

=====
Train on 35000 samples, validate on 15000 samples

Epoch 1/35

35000/35000 [=====] - 11s 303us/sample - loss: 1.6156 - accuracy: 0.4233

Epoch 2/35

35000/35000 [=====] - 8s 237us/sample - loss: 1.2252 - accuracy: 0.5632

Epoch 3/35

35000/35000 [=====] - 8s 233us/sample - loss: 1.0634 - accuracy: 0.6223

Epoch 4/35

35000/35000 [=====] - 8s 237us/sample - loss: 0.9754 - accuracy: 0.6570

Epoch 5/35

35000/35000 [=====] - 8s 232us/sample - loss: 0.9038 - accuracy: 0.6839
 Epoch 6/35
 35000/35000 [=====] - 8s 232us/sample - loss: 0.8506 - accuracy: 0.6985
 Epoch 7/35
 35000/35000 [=====] - 8s 231us/sample - loss: 0.8043 - accuracy: 0.7163
 Epoch 8/35
 35000/35000 [=====] - 8s 230us/sample - loss: 0.7660 - accuracy: 0.7327
 Epoch 9/35
 35000/35000 [=====] - 8s 230us/sample - loss: 0.7400 - accuracy: 0.7378
 Epoch 10/35
 35000/35000 [=====] - 8s 229us/sample - loss: 0.7072 - accuracy: 0.7497
 Epoch 11/35
 35000/35000 [=====] - 8s 229us/sample - loss: 0.6877 - accuracy: 0.7552
 Epoch 12/35
 35000/35000 [=====] - 8s 231us/sample - loss: 0.6605 - accuracy: 0.7678
 Epoch 13/35
 35000/35000 [=====] - 8s 238us/sample - loss: 0.6458 - accuracy: 0.7737
 Epoch 14/35
 35000/35000 [=====] - 8s 236us/sample - loss: 0.6313 - accuracy: 0.7766
 Epoch 15/35
 35000/35000 [=====] - 8s 232us/sample - loss: 0.6134 - accuracy: 0.7813
 Epoch 16/35
 35000/35000 [=====] - 8s 234us/sample - loss: 0.5958 - accuracy: 0.7891
 Epoch 17/35
 35000/35000 [=====] - 8s 232us/sample - loss: 0.5828 - accuracy: 0.7956
 Epoch 18/35
 35000/35000 [=====] - 8s 242us/sample - loss: 0.5747 - accuracy: 0.7988
 Epoch 19/35
 35000/35000 [=====] - 9s 244us/sample - loss: 0.5575 - accuracy: 0.8028
 Epoch 20/35
 35000/35000 [=====] - 8s 232us/sample - loss: 0.5447 - accuracy: 0.8070
 Epoch 21/35
 35000/35000 [=====] - 8s 232us/sample - loss: 0.5302 - accuracy: 0.8095
 Epoch 22/35
 35000/35000 [=====] - 8s 229us/sample - loss: 0.5255 - accuracy: 0.8133
 Epoch 23/35
 35000/35000 [=====] - 8s 234us/sample - loss: 0.5174 - accuracy: 0.8172
 Epoch 24/35
 35000/35000 [=====] - 8s 231us/sample - loss: 0.5120 - accuracy: 0.8191
 Epoch 25/35
 35000/35000 [=====] - 8s 235us/sample - loss: 0.4956 - accuracy: 0.8251
 Epoch 26/35
 35000/35000 [=====] - 8s 231us/sample - loss: 0.4912 - accuracy: 0.8255
 Epoch 27/35
 35000/35000 [=====] - 8s 228us/sample - loss: 0.4824 - accuracy: 0.8304
 Epoch 28/35
 35000/35000 [=====] - 8s 227us/sample - loss: 0.4786 - accuracy: 0.8303
 Epoch 29/35
 35000/35000 [=====] - 8s 225us/sample - loss: 0.4684 - accuracy: 0.8327

```

Epoch 30/35
35000/35000 [=====] - 8s 229us/sample - loss: 0.4626 - accuracy: 0.8383
Epoch 31/35
35000/35000 [=====] - 8s 236us/sample - loss: 0.4549 - accuracy: 0.8368
Epoch 32/35
35000/35000 [=====] - 8s 231us/sample - loss: 0.4508 - accuracy: 0.8383
Epoch 33/35
35000/35000 [=====] - 8s 233us/sample - loss: 0.4411 - accuracy: 0.8430
Epoch 34/35
35000/35000 [=====] - 8s 233us/sample - loss: 0.4355 - accuracy: 0.8463
Epoch 35/35
35000/35000 [=====] - 8s 237us/sample - loss: 0.4394 - accuracy: 0.8428

```

```
In [17]: test_loss, test_acc = model1.evaluate(x_test, y_test, verbose = 0)
```


```
print('\nTest accuracy:', test_acc)
```

Test accuracy: 0.7944

```
In [18]: #plot
plt.figure(figsize=(10,9))

plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('accuracy')
plt.legend(['train', 'test'])
plt.grid()
plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('loss')
plt.legend(['train', 'test'])
plt.grid()

plt.show()
```



Usar red neuronal para reconocer objetos

```
In [19]: x_test[[0]].shape
```

```
Out[19]: (1, 32, 32, 3)
```


```
In [20]: np.argmax( model1.predict(x_test[[0]]) )
```

```
Out[20]: 3
```

```
In [21]: plt.imshow(x_test[0])
plt.axis('off')
plt.title( 'Etiqueta: ' + dict[ int(y_test[0])] + ' '
          'Prediccion: ' + dict[ np.argmax( model1.predict(x_test[[0]]) ) ] )
```

```
Out[21]: Text(0.5, 1.0, 'Etiqueta: Cat Prediccion: Cat')
```



Etiqueta: Cat Prediccion: Cat



```
In [22]: plt.imshow(x_test[1])  
         plt.axis('off')  
         plt.title( 'Etiqueta: ' + dict[ int(y_test[1])] + ' '  
                   'Prediccion: ' + dict[ np.argmax( model1.predict(x_test[[1]]) ) ] )
```

```
Out[22]: Text(0.5, 1.0, 'Etiqueta: Ship Prediccion: Ship')
```

Etiqueta: Ship Prediccion: Ship



- Experimentar con la arquitectura de la red, agregando neuronas y capas
- Agregar y modificar la regularización para entrenar en menor tiempo
- Mejorar test accuracy